



UNIVERSITÀ
DI TRENTO

Department of Information Engineering and Computer Science

Master's Degree in
Computer Science

FINAL DISSERTATION

A POLICY-DRIVEN KUBERNETES-BASED
ARCHITECTURE FOR RESOURCE
MANAGEMENT IN MULTI-CLOUD
ENVIRONMENTS

Supervisor

Prof. Sandro Luigi Fiore

Student

Leonardo Vicentini

Co-supervisors

Dott. Diego Braga

Dott. Francesco Lumpp

Academic year 2023/2024

Acknowledgements

Thanks to my Family and Friends for the support and encouragement throughout the years.

Contents

Abstract	4
1 Introduction	5
1.1 Context	5
1.2 Problem statement	5
1.3 Method	6
1.4 Personal contribution	6
2 Background	7
2.1 Public cloud providers	7
2.1.1 Cloud Regions and Availability Zones	7
2.1.2 Multi-cloud paradigm	8
2.2 Kubernetes	9
2.2.1 Kubernetes extendability	11
2.2.2 Kubernetes as a platform	11
2.2.3 Helm	11
2.3 Krateo PlatformOps	12
2.3.1 Krateo core-provider	12
2.3.2 Krateo composition-dynamic-controller	12
2.4 Multi cloud resource management	13
2.4.1 Dynamic Virtual Machine placement	13
2.4.2 Cloud service brokers	13
2.4.3 AI-based resource management in cloud computing	14
2.4.4 Policy-driven resource management systems	14
2.5 GreenOps landscape	15
2.5.1 Green Software foundation	15
2.5.2 Computational Sustainability by Public cloud providers	16
2.6 Carbon-aware systems for resource management	17
2.6.1 CASPER	17
2.6.2 CASPIAN	17
2.6.3 CarbonScaler	17
2.6.4 A Low Carbon Kubernetes Scheduler	18
2.6.5 Microsoft’s Carbon-Aware Kubernetes strategy	18
2.7 Multi-cloud resource management - major takeaways	18
3 System design and implementation	19
3.1 Assumptions	19
3.1.1 Workload definition	19
3.1.2 System limitations	20
3.2 System Architecture	21
3.3 Krateo PlatformOps integration	21
3.3.1 Resource management: the Custom Kubernetes “Synchronization Operator” approach	22

3.3.2	Resource management: the Krateo PlatformOps approach	23
3.4	Multi-Cloud Integration through Kubernetes Operators	29
3.4.1	Azure Kubernetes Operator	29
3.4.2	GCP Operator	30
3.4.3	AWS Operator	31
3.5	Kubernetes Mutating Webhook Configuration	34
3.6	Open Policy Agent (OPA)	35
3.6.1	Policy as Code paradigm	35
3.6.2	OPA architecture overview	36
3.6.3	OPA and external data sources	37
3.6.4	OPA integration with Kubernetes	38
3.6.5	OPA policies	39
3.6.6	OPA policy bundles	40
3.6.7	Latency policy	41
3.6.8	GDPR policy	43
3.6.9	Scheduling outcome policy	43
3.6.10	OPA Data mapping	44
3.6.11	OPA integration in the system	45
3.6.12	OPA Gatekeeper	46
3.6.13	OPA advanced features	48
3.7	MLOps infrastructure	49
3.7.1	MLOps purpose	49
3.7.2	MLflow	49
3.7.3	KServe	51
3.7.4	MLOps general architecture	53
3.8	Cloud resource metrics	54
3.8.1	System performance metrics	54
3.8.2	Power consumption metrics	55
3.8.3	Carbon metrics	56
3.8.4	Impact framework potential integration	56
3.9	End-to-End workflow	57
4	Conclusion	59
4.1	End-to-end integrated test	59
4.2	GreenOps system evaluation	59
4.2.1	Theoretic upper bound	59
4.2.2	Baseline definition	59
4.2.3	Black hole phenomenon	59
4.2.4	Side effects	59
4.2.5	Preliminary evaluation	59
4.3	Future improvements	59
4.3.1	Multi model serving	60

List of Figures

1.1	Project parts	5
2.1	Worldwide market share of leading cloud infrastructure service providers as of Q3 2024 [?]	8
2.2	Geo-distribution of Azure data centers with country Grid Carbon Intensity	9
2.3	Kubernetes architecture by CNCF [?]	10
2.4	Operator paradigm	11
2.5	Krateo core-provider and composition-dynamic-controller architecture [?]	13
3.1	Multi-cloud resource management with Custom Kubernetes “Synchronization Operator” approach	23
3.2	Multi-cloud resource management with Krateo PlatformOps approach	24
3.3	Minimum set of Azure resources for VM provisioning	30
3.4	Minimum set of GCP resources for VM provisioning	31
3.5	Minimum set of AWS resources for VM provisioning	32
3.6	Kubernetes Admission Control	34
3.7	Kubernetes Mutating Webhook example [?]	35
3.8	OPA architecture	37
3.9	Kubernetes mutating webhook and OPA integration	38
3.10	OPA policy bundles	41
3.11	Latency matrix example (Azure regions subset)	42
3.12	OPA Data mapping	45
3.13	General architecture	46
3.14	MLflow deployment configuration	50
3.15	Model deployment configuration and GreenOps system integration	52
3.16	MLOps Architecture	53
3.17	Standard Prometheus architecture for system performance metrics collection	55
3.18	Sequence diagram of the end-to-end workflow	58

Abstract

contesto

motivazioni

riassunto problema affrontato

tecniche utilizzate analisi requisiti, analisi pprogetti/prodotti disponibili creazione proof of concept

risultati raggiunti: e2e testing

contributo personale

1 Introduction

[TO BE ADDED] The work described in this thesis is part of a larger project that aims to develop a system ... The project is divided into three main parts: data analysis, machine learning, and cloud infrastructure as shown in Figure 1.1.

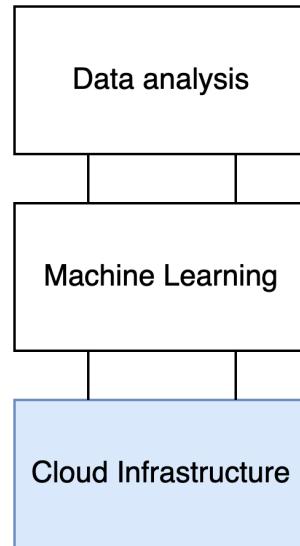


Figure 1.1: Project parts

1.1 Context

[TO BE ADDED]

MULTI Cloud MULTI CLOUD RESOURCE MANAGEMENT

Computational sustainability

GreenOps

GreenOps for FinOps (Operating for GreenOps may lead to reduced costs)

Geographical shifting and Time shifting Carbon-aware workload scheduling psrticular kind of workload (delay tolerant)

Cloud sustainability

Current Sustainable Cloud Computing Landscape we are in the infrastructure tooling section in particular scheduling (day 1 operatSION)

scaling and resource tuning are usuaually day 2 operation

the system was envisioned with this in mind and is capable of doing that

reducing the Carbon Footprint (CFP) due to executing workloads, while satisfying Quality of Service (QoS) requirements

1.2 Problem statement

[TO BE ADDED]

Use cases (basic ones for the beginning) higher level explanation here first use case ("GreenOps" VM scheduling)

second: scaling down a vm infrastructure already put in place

the system was designed with flexibility in mind therefore a workload could be potentially anything the condition is just to be represented in some way and have something else do certain actions based

on that representation As we will see in section XXX, the most simple of this would be K8s operators this is described in section XYZ

1.3 Method

[TO BE ADDED]

Developing a real solution, integrating it on top of OSS production-ready solution

we are on the consumer side, not on the provider side

System architecture to start with: Saima's + Krateo platform integration into an existing platform (krateo) leveraging krateo components

Krateo Core Provider and cdc instead of developing 1 or more K8s operators from scratch analysis of possible solutions implemented poc

Initial analysis of a solution with operators were tried

A PoC comprising 1 operator was created “Synchronization operation” cons: maintainer costs ideation and creation of architectural diagrams

tackling first use case but create a system that is flexible enough to be used for other use cases as well

1.4 Personal contribution

[TO BE ADDED]

The project, ideated and supervised by Prof. Fiore is mainly divided into 3 parts.

exploratory data analysis data preparation

model training model selection

infrastructure part

GOAL The goal of the project is to employ mainly time-shifting and geographical shifting for the scheduling of workload leveraging a multi-cloud setting. We can choose to schedule workloads in periods and regions with low carbon intensity (when renewables are plentiful). Therefore, targeted workloads are the ones that are not time-sensitive but instead are quite delay-tolerant. For example training a machine learning model could wait until a period of low carbon intensity. Another example is shifting video / image processing, as Google is doing. Kubernetes is leveraged as a platform for scheduling and managing workloads on different cloud providers. Long term goal: “Using electricity when the carbon intensity is low is the best way to ensure investment flows towards low-carbon emitting plants and away from high-carbon emitting plants”.

2 Background

In this chapter, we provide an overview of the main concepts, paradigms and technologies that are relevant for the purpose of this work. We start by introducing the concept of **Public Cloud Providers** and the **Multi Cloud paradigm**. We then provide a brief overview of **Kubernetes**, in particular focusing on the concept of **Kubernetes as a platform** and the **Helm** package manager. After that, **Krateo PlatformOps**, an open-source Kubernetes-based platform that is a fundamental part of our system, is introduced. We deemed also valuable to provide an overview of the existing works in the field of **multi-cloud resource management** to highlight recurrent patterns and design choices. Finally, the **GreenOps landscape** is introduced, focusing on the the **Computational Sustainability** initiatives by Public Cloud Providers and on **carbon-aware systems for resource management**.

2.1 Public cloud providers

The Cloud Computing definition by the National Institute of Standards and Technology (NIST) [?] states that “*Cloud Computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction*” [?]. **Public Cloud Providers** or Cloud Service Providers (CSPs) are companies that have as their core business the provisioning of cloud computing services. These services, which are growing in number and complexity, range from computing resources to storage, networking, databases, machine learning, and more. Public cloud is a deployment model that allows organizations to consume cloud services without having to build and maintain their own physical infrastructure (e.g., private cloud), therefore reducing capital expenditure (CapEx) and shifting to an operational expense model (OpEx). The public cloud model is actually a multi-tenant environment where physical resources maintained and operated by the cloud provider are shared among multiple customers (tenants). These resources are offered in the form of various services at different levels of abstraction (e.g., Infrastructure as a Service, Platform as a Service, Software as a Service) and are provided on-demand on a pay-as-you-go basis. Currently, as of 2024, the public cloud market is dominated by three main hyperscalers: **Amazon Web Services (AWS)**, **Microsoft Azure**, and **Google Cloud Platform (GCP)**. An hyperscaler is a company (cloud service provider) that operates a data center infrastructure at a massive scale and is able to provide cloud services to a global audience. Figure 2.1 shows the worldwide market share of leading cloud infrastructure service providers as of Q3 2024 [?].

2.1.1 Cloud Regions and Availability Zones

With the term **cloud region**, cloud providers refer to a geographical area where they have one or more data centers. Cloud providers usually further divide regions into availability zones which main purpose is to provide high availability and fault tolerance. For the purpose of this work, we will consider the concept of **cloud region** as the primary unit of deployment of cloud resources. Usually, by cloud providers design choice, each cloud region supports a subset of the available cloud services and instance types. We could safely assume that our targeted workload specifications are quite standard and therefore can be scheduled on any cloud region. As of late 2024, the three major cloud service providers have established extensive global infrastructures. Table 2.1 shows the number of regions and availability zones (AZs) for each provider as of Q4 2024 - Q1 2025 [?] [?]. We must note that the number of regions and availability zones is constantly changing as cloud providers are expanding their global infrastructure.

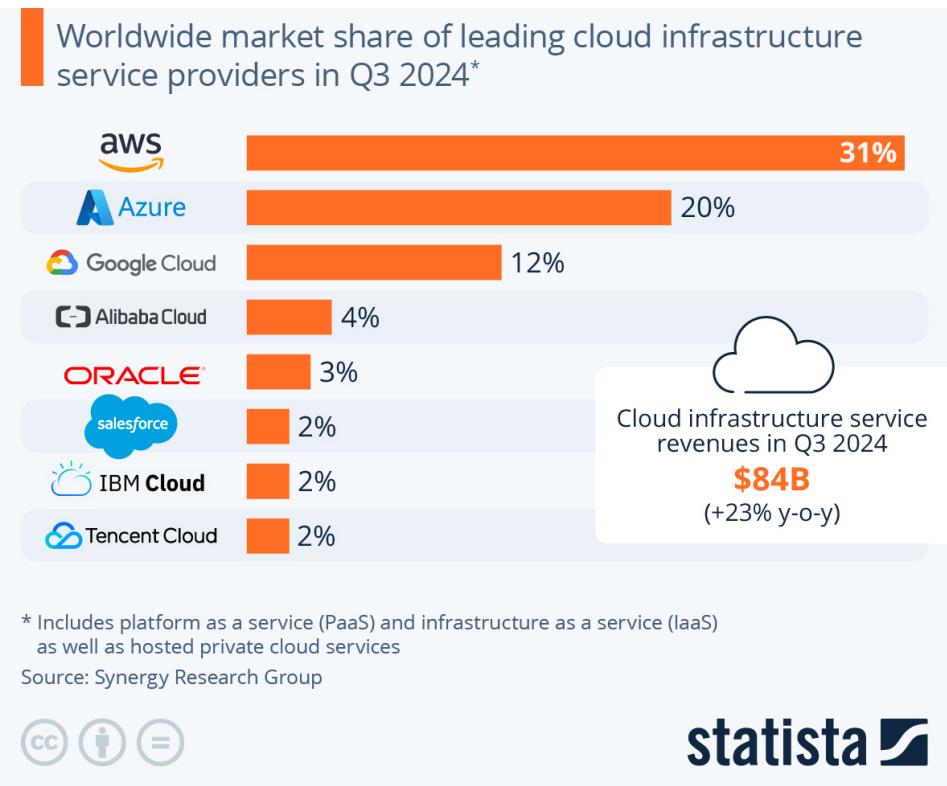


Figure 2.1: Worldwide market share of leading cloud infrastructure service providers as of Q3 2024 [?]

Cloud Provider	Regions	Availability Zones
Amazon Web Services (AWS)	36	114
Microsoft Azure	33	93
Google Cloud Platform (GCP)	40	121

Table 2.1: Number of Cloud Regions and Availability Zones by Provider as of 2024 [?]

We must also note that each cloud provider has a different number of regions and zones and also different naming conventions for them. There is no standardization nor a convention in the industry for this and cloud providers have chosen their own naming schemes. For instance a region located in London is called “eu-west-2” in AWS, “uk-west” in Azure and “europe-west2” in GCP. This detail must be taken into account when designing and implementing a multi-cloud system. For what concerns the geo-distribution of cloud regions, as an example, Figure 2.2 shows the location of **Azure cloud data centers** around the world and the country Grid Carbon Intensity as reported by Electricity Maps (year 2024) [?]. Data center coordinates are retrieved by an Azure CLI command dump [?].

2.1.2 Multi-cloud paradigm

The multi-cloud paradigm refers to the strategic utilization of cloud services from multiple public cloud providers within a single, heterogeneous architecture. This approach allows organizations to distribute workloads, applications, and data across multiple public cloud providers [?]. On the other hand, the hybrid cloud paradigm refers to utilization of both public cloud services and on-premises infrastructure (private cloud). The multi-cloud strategy is often adopted by organizations to **avoid vendor lock-in, increase redundancy, increase flexibility and optimize costs** [?]. For what concern flexibility and optimization, a multi-cloud strategy allows an organization to select the best service on a case-by-case basis, leveraging the strengths of each provider. It could even be the case that, for the end user, the choice of the cloud provider for a service is transparent meaning that the organization itself (or the system) is the one that decides which cloud provider to use for a specific service based on some criteria. Vendor lock-in is avoided by designing applications to be **cloud-agnostic** and by leveraging open-source technologies (e.g., Kubernetes) and standards [?]. If

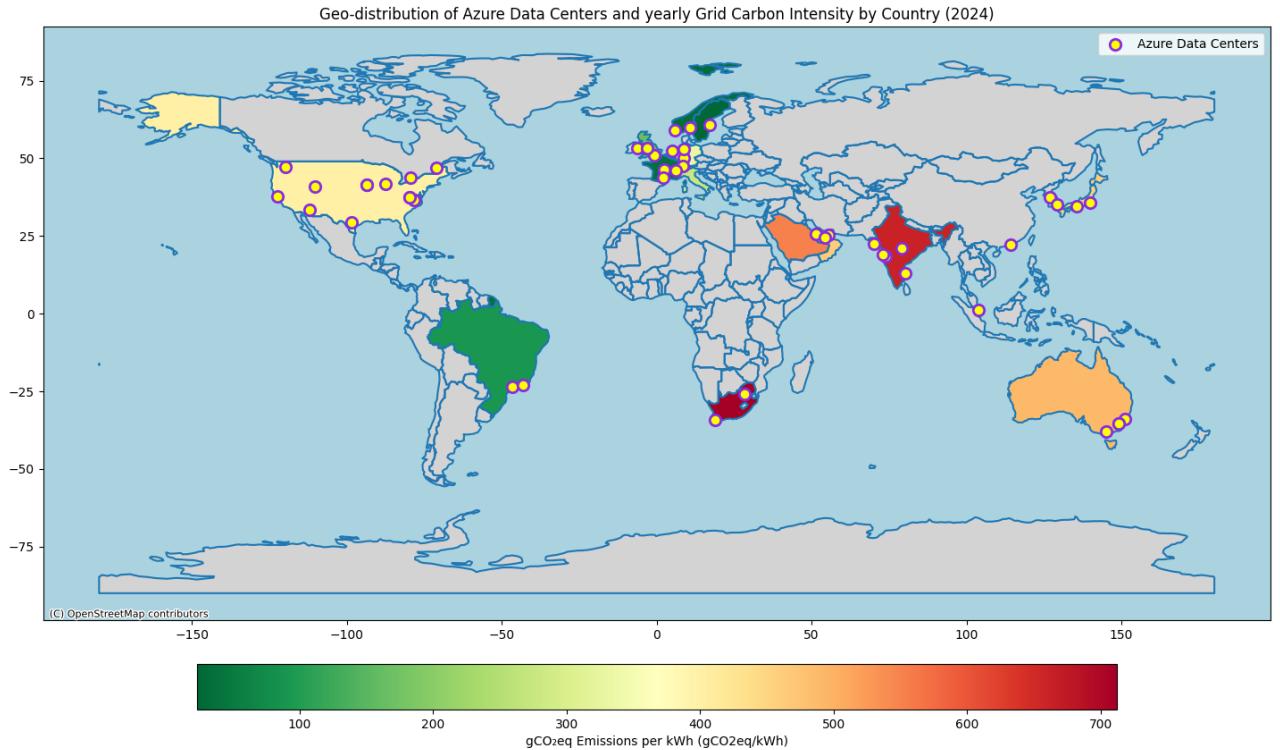


Figure 2.2: Geo-distribution of Azure data centers with country Grid Carbon Intensity

an application is designed to be cloud-agnostic, it can be relatively easily migrated from one cloud provider to another. Multi-cloud is a strategy that might be more difficult to implement compared to a single cloud strategy but the long-term benefits are significant. Challenges and disadvantages are also present in a multi-cloud strategy [?]. As a matter of fact it can be safely said that **complexity in the overall infrastructure management is increased**. Another challenge is the **integration of services** across different cloud providers. An additional consideration could be done in terms of security: a multi-cloud strategy can **increase the attack surface** of an organization, therefore security measures must be carefully implemented.

For the purpose of this work, adopting a multi-cloud strategy is beneficial for several reasons. First and foremost, **user-centric flexibility** is achieved bringing the benefits described above. If a user or organization has a preference for a specific cloud provider, the system can be configured to use only that provider or a specific subset of providers. Secondly, the system can be designed to be **cloud-agnostic**, meaning that the choice of the cloud provider is transparent to the end user. With some specific configurations, the subset of providers to be used can be decided by the system itself based on some criteria. Currently, as described in section 3.6.9, the user can specify a subset of cloud providers (among the 3 major ones) to be used for the deployment of resources. Finally, in theory, since different cloud providers have data centers in various locations around the world, some low-carbon regions might be available only on a specific cloud provider.

2.2 Kubernetes

Kubernetes is an open-source platform for automating the orchestration of containerized applications. It is widely used in the industry and became the **de-facto standard for container orchestration**. An extensive description of Kubernetes is out of the scope of this thesis but it is deemed necessary to provide a brief overview of the main concepts of Kubernetes that are relevant for the purpose of this work. Figure 2.3 shows the Kubernetes architecture as described by the Cloud Native Computing

Foundation (CNCF) [?].

Kubernetes architecture

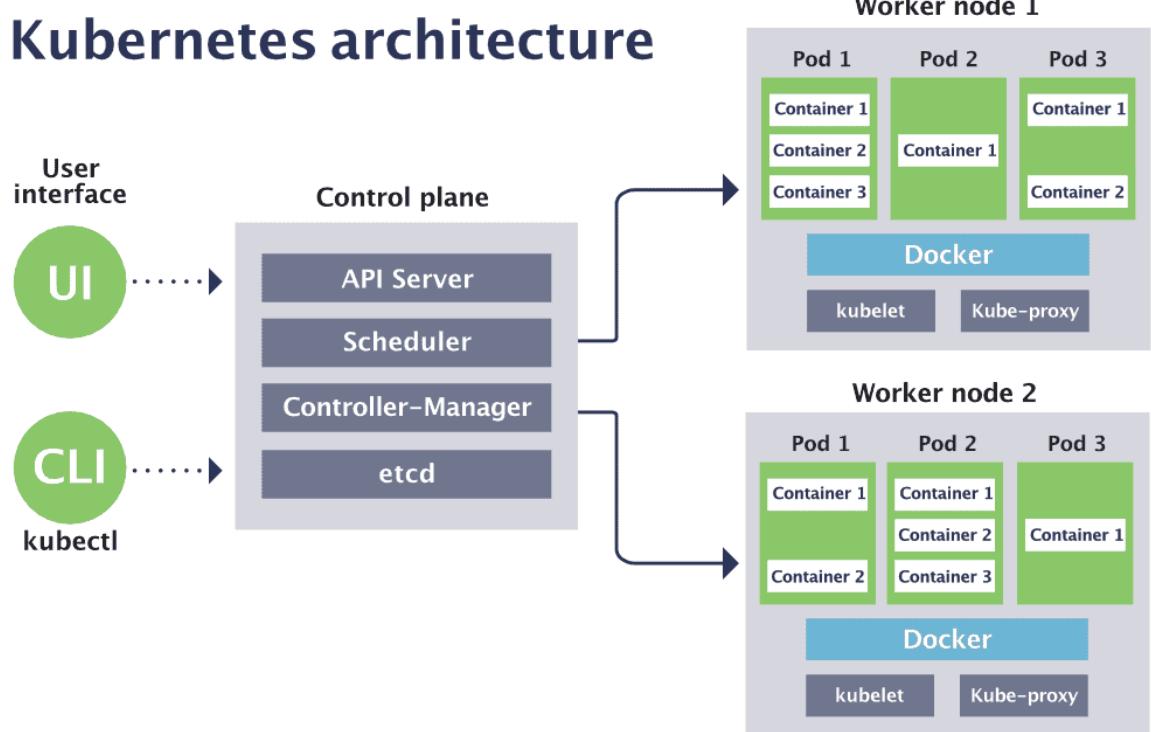


Figure 2.3: Kubernetes architecture by CNCF [?]

The main components of Kubernetes are:

- **Kubernetes API server:** the central component that manages the Kubernetes cluster. It exposes the Kubernetes API and is responsible for validating and mutating data related to Kubernetes objects.
- **etcd:** a key-value store used to store the cluster state.
- **Kubernetes controller manager:** a daemon that embeds the core control loops shipped with Kubernetes.
- **Kubernetes scheduler:** a component that assigns Pods to nodes.
- **Kubelet:** an agent that runs on each node in the cluster. It makes sure that containers are correctly running in a Pod.
- **Kubernetes proxy:** a network proxy that runs on each node in the cluster. It maintains network rules and it is responsible for routing network traffic.
- **Container runtime:** the software that is responsible for running containers on the Node (in this case represented by Docker).

It must be noted that in our system we are not extending the Kubernetes scheduler as described for instance in the works described in section 2.6.4 and 2.6.5, since we are not dealing with the scheduling of in-cluster resources (e.g., Kubernetes Pods) nor with the provisioning of entire Kubernetes clusters. We are instead focusing on the **management of external resources on cloud providers** (only VMs in this first iteration), leveraging Kubernetes as a control plane for the management of these resources. As described in the following section, our focus will be on the Kubernetes API server and in particular on Kubernetes Admission Control.

2.2.1 Kubernetes extensibility

Kubernetes allows for the extension of its functionalities through the use of **Custom Resource Definitions (CRDs)** and **Kubernetes Operators**, effectively adopting the so-called **Operator paradigm**. Simply put, Custom Resource Definitions are a way to instruct Kubernetes to manage new resource types. They are a **schema** that defines the structure of the resource and the Kubernetes API server will validate the resource against the schema. **Custom Resources (CRs)** are actually instances of the resources defined by CRDs and are managed by an Operator. The Operator is a piece of software (controller) that is responsible for managing the lifecycle of the resources defined by the CRD. Effectively the code of an Operator is usually deployed on the Kubernetes cluster in the form of a Deployment. This concept is not much different of what actually happens for standard built-in Kubernetes resources which are however managed by built-in controllers (e.g., Deployment controller, ReplicaSet controller, part of Kubernetes controller manager). An high-level overview of the Operator paradigm is depicted in Figure 2.4.

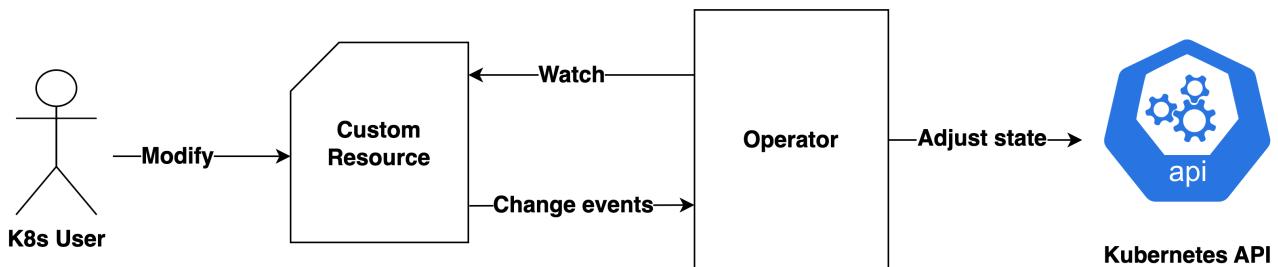


Figure 2.4: Operator paradigm

2.2.2 Kubernetes as a platform

Recently, the paradigm of leveraging Kubernetes as a platform to manage external resources has become more and more popular. Therefore, Kubernetes use cases have expanded beyond the management of containerized applications to the management of any desired resource on any infrastructure. The idea is to **represent external resources as Kubernetes objects** and leverage Kubernetes' declarative paradigm and robust API to achieve unified and efficient infrastructure management. This approach can be exemplified by the public cloud providers' operators which extend Kubernetes to manage their cloud resources as described in section ??.

Config Connector is a Kubernetes operator developed by Google Cloud that allows to configure numerous Google Cloud services and resources using Kubernetes tooling and APIs [?]. The idea general idea is that organizations might have to deal with a **heterogeneous infrastructure** composed of different cloud providers and on-premises resources. By **unifying infrastructure management under Kubernetes**, complexity is reduced and velocity is increased because the same tools and APIs can be used to manage all the resources in a consistent way [?].

2.2.3 Helm

Helm is a **package manager** for Kubernetes. Therefore with Helm it is possible to define, install and manage Kubernetes applications in a simpler way compared to a manual management of Kubernetes resources manifests [?]. Helm is a graduated project in the CNCF and it is the de-facto standard for Kubernetes package management. The key concept is the **Helm chart**, which is a collection of files that describe a related set of Kubernetes resources. These files are mainly of two types: templates and values. The **templates** are Kubernetes manifest files that are rendered by Helm's **powerful templating engine**. The **values** are the set of variables that are used to render the templates. Upon an Helm chart installation, Helm will render the templates "injecting" the values and deploy the resources in the Kubernetes cluster. One major advantage that Helm provides is the complete management of the lifecycle of the resources. As a matter of fact, Helm allows to easily **upgrade**, **rollback** and **uninstall** the Kubernetes resources deployed with a Helm chart reducing time and errors in such operations [?]. Without Helm, the user would have to deal with each single Kubernetes resource manifest file and manually apply changes to them. Finally, users can benefit of Helm charts

already developed by the community and leverage chart distribution within their organization using Helm repositories.

2.3 Krateo PlatformOps

Krateo PlatformOps (Krateo) is an **open-source Kubernetes-based platform** that aims to provide a unified interface for managing any desired resource on any infrastructure [?]. Krateo runs as a Kubernetes deployment inside a Kubernetes cluster but **acts as a control plane** even for resource external to the Kubernetes cluster. The only requirement for this management is that the resources need to be logically describable using a YAML file which represents the desired state of the resources [?]. Krateo is composed of three main parts:

- Krateo Composable Operations
- Krateo Composable Portal
- Krateo Composable FinOps

For the purpose of this work, we will focus on the **Krateo Composable Operations** part, which is the core of the Krateo platform and is responsible for managing the lifecycle of resources in a Kubernetes cluster [?]. Krateo Composable Operations is composed in turn by several components. Due to their core importance in our system, we will briefly describe the functionalities of the **Krateo core-provider** and the **Krateo composition-dynamic-controller** as described in Krateo official documentation [?] and [?].

2.3.1 Krateo core-provider

The Krateo core-provider, as its name suggests, is the core component of the Krateo platform. It is a **Kubernetes operator** that has the duty of downloading and managing Helm charts. It first checks for the existence of a file named *values.schema.json* in the chart folder and uses it to generate a Kubernetes Custom Resource Definition (CRD), accurately representing the possible values that can be expressed for the installation of the chart [?]. The file *values.schema.json* is a JSON schema that describes the structure of the *values.yaml* file for the related Helm chart and it is considered a standard best practice for Helm charts. It basically provides a way to validate the *values.yaml* file before the Helm chart is installed (i.e., to check if the values are in the correct format and if all the required values are present) [?]. Therefore the file *values.schema.json* could be useful in the context of DevOps practices and CI/CD pipelines but in the case of Krateo it is pivotal for the generation of the CRD. In other words, the Krateo core-provider operator is responsible for deploying the Helm chart as a **native Kubernetes resource**, which allows for the management of the whole Helm chart lifecycle through Kubernetes APIs [?]. As a matter of fact, out of the box, Kubernetes does not provide a way to manage Helm charts natively and the Krateo core-provider is one tool that allows to do so. The Kubernetes Custom Resource Definition introduced by the Krateo core-provider is called **CompositionDefinition**. It is a CRD that represents the Helm chart and its values (a Helm Chart archive (.tgz) with a JSON Schema for the *values.yaml* file) [?]. Upon a CompositionDefinition manifest application to the Kubernetes cluster, the Krateo core-provider generates the CRD from the schema defined in the *values.schema.json* file included in the chart. It then deploys an instance of the Krateo composition-dynamic-controller, setting it up to manage resources of the type defined by the CRD [?]. Therefore, this is an example of the Operator paradigm described in section 2.2.1 where the Krateo core-provider is the Operator and the CompositionDefinition is the Custom Resource Definition.

2.3.2 Krateo composition-dynamic-controller

The Krateo composition-dynamic-controller is the Kubernetes operator that is instantiated by the Krateo core-provider to manage the specific Custom Resources whose Custom Resource Definition is generated by the core-provider. In practice, when a Custom Resource (CR) is created, the specific instance of composition-dynamic-controller checks if a Helm release associated with the CR already exists in the cluster [?]. If this is not the case, it performs an *helm install* operation using the values specified in the CR to create a new Helm release. This will practically deploy all the resources defined in the Helm chart using **Helm's templating engine**. This feature is particularly important in the

context of the system described in this thesis and in particular for multi-cloud resource management as explained in section XYZ. However, if the Helm release does already exist, it instead executes an *helm upgrade* operation, updating the release's values with those specified in the CR, effectively updating the resources in the cluster. This mechanism will be leveraged as well for a particular feature of our system described in section XYZ. Finally, when the CR is deleted from the cluster, the instance of the composition-dynamic-controller performs an *helm uninstall* on the release, removing all the resources defined in the Helm chart from the cluster [?].

It must be said that Krateo composition-dynamic-controller has a set of additional features that are not described in this work such as multi-version support. The architecture of the Krateo core-provider and Krateo composition-dynamic-controller is depicted in Figure 2.5, as described and depicted in [?].

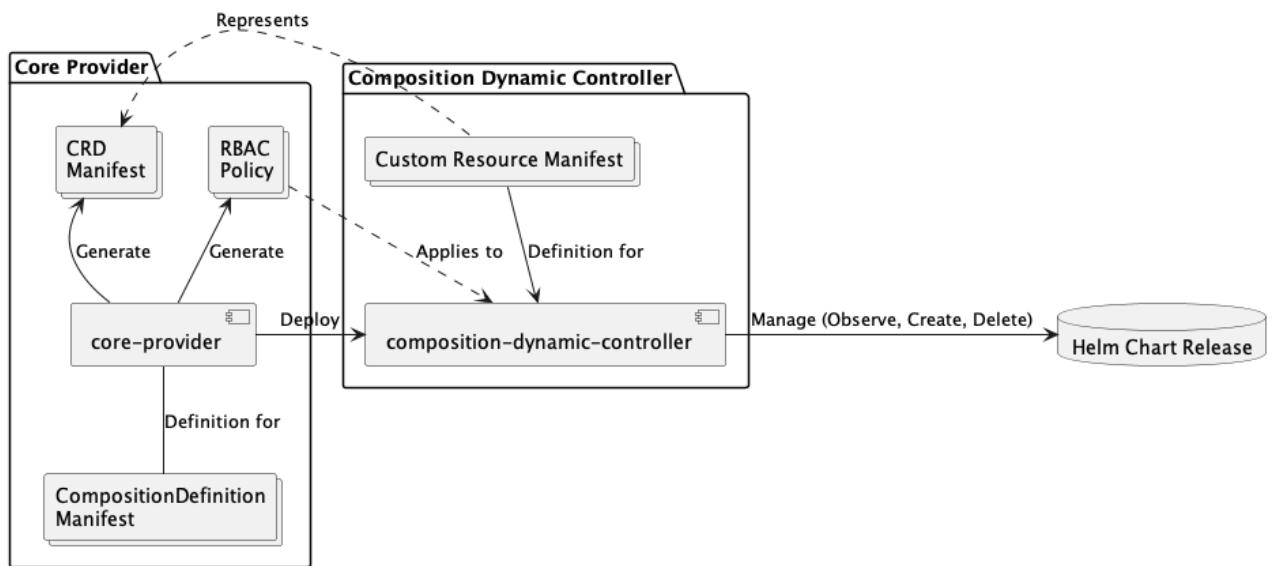


Figure 2.5: Krateo core-provider and composition-dynamic-controller architecture [?]

2.4 Multi cloud resource management

The idea of a **dynamic management of workloads leveraging a multi-cloud paradigm** is not new. In this section we will provide an overview of some of the existing works in the literature that have tackled the problem of multi-cloud resource management.

2.4.1 Dynamic Virtual Machine placement

The work of Simarro et al. [?] back at the dawn of cloud computing (2011) proposed a multi-cloud architecture for the dynamic placement of Virtual Machines (VMs). The main objective of the system was cost optimization but this paradigm provides reliability and flexibility as well. The scheduling part is comprised of a “**cloud broker**” that is responsible for VM placement **transparent to users** providing a single uniform interface to the cloud resources. Users can provide to the system a “**service description template**” to specify the number of VMs to provision and some constraints. The cloud broker architecture is composed of two major components: the **scheduler** and the **cloud manager** [?]. The former is responsible for placement decision across multiple cloud providers, while the latter is responsible for the actual management of the VMs in the cloud providers. More precisely, the cloud manager is represented by the OpenNebula (ONE) virtual infrastructure manager [?]. OpenNebula is an open-source platform that aims to provide a unified management interface for multiple virtualization technologies and cloud providers [?].

2.4.2 Cloud service brokers

A CSB is a system that acts as an **intermediary** between cloud service providers and consumers, providing a **unified interface** to manage cloud resources across multiple providers [?]. Cloud service brokers (CSBs) were described and categorized by Wadhwa et al. [?] in their work of 2013. The

emerging market of cloud computing led to the proliferation of cloud services and providers, and by consequence the need for mechanisms to manage costs, capacity and resources [?].

An interesting CSB example in the literature is the **STRATOS** system by Pawluk et al. proposed in 2012. The work can be considered a pioneer in the field of multi-cloud resource management since it can be framed in the first years of cloud computing [?] but the proposed paradigms and concepts are relevant today. STRATOS tried to **avoid the assumption of resource homogeneity** and represented an initial attempt to provide a “**cross-cloud resource provisioning**” system [?]. The proposed architecture enables the specification of high-level objectives that can be assessed in a standardized manner across different providers. The decision-making process is fully automated, shifting the decision point from deployment to runtime [?]. Users first submit a topology document, triggering the Cloud Manager to communicate with the Broker for topology instantiation. The Broker (implemented in Java) then conducts the initial resource acquisition decision (RAD), optimizing the allocation of resources across multiple providers (configured beforehand) [?]. Experiments indicate that distributing workloads across different cloud providers can reduce the overall cost of the topology. The approach taken by the authors primarily focuses on two objectives: **cost efficiency** and **avoiding vendor lock-in**. The application environment was deployed on public cloud platforms, specifically AWS and Rackspace [?].

[TO BE ADDED] A Kubernetes ‘Bridge’ operator between cloud and external resources

2.4.3 AI-based resource management in cloud computing

More recent works have focused on the development of systems that leverage AI techniques for the optimization of resource management in cloud computing environments.

The work of 2022 by Khan et al. [?] provides a comprehensive review of the state of the art in the field of machine learning (ML)-centric resource management in cloud computing. Although this work focuses on the cloud provider side (**resource management in data centers**), it provides some interesting insights that can be leveraged also at different levels of the cloud computing stack. The authors highlight the fact that traditionally, only static policies were used for resource management in cloud computing, but the advent of ML techniques has enabled the development of more dynamic and adaptive resource management systems [?].

Tuli et al. in their work of 2021 [?] propose a system named **HUNTER** which stands for “Holistic resoUrce maNagemenT technique for Energy-efficient cloud computing using aRtificial intelligence”. In particular, in this work, tasks are modeled as **containers instances**. For the purpose of this thesis, we can describe how this system was implemented from a infrastructural point of view. The system is composed of four main components [?]:

- **Cloud Workload Management Portal:** a web-based portal that allows users to submit their workloads and specify the requirements (SLA constraints, QoS constraints, etc.).
- **Workload Manager:** a component that is responsible for the processing of incoming workloads.
- **Cloud Broker:** the component that is responsible for the allocation of resources to cloud worker nodes
- **Cloud Hosts:** set of cloud worker nodes (in the form of both private and public cloud)

The Cloud Broker is the core of the system and is further divided into three sub-managers: Service Manager, CDC Manager, and Resource Manager. The Service Manager is responsible for SLAs and QoS constraints, the CDC Manager is responsible for the actual allocation (provisioning) and migration of resources to cloud worker nodes and resource monitoring and the Resource Manager is responsible for the scheduling of tasks and contains the actual sustainability models (energy, thermal, cooling.) [?].

2.4.4 Policy-driven resource management systems

García García et al. (2014) propose **Cloudcompaas** [?], a Service Level Agreement (SLA)-driven platform for dynamic cloud resource management, focusing on the automation of resource provision-

ing, scheduling, and monitoring. The aim of the work is to provide a SLA-aware Platform-as-a-Service platform for the entire management of cloud resource lifecycle [?]. The targeted resource to be managed spans from IaaS to PaaS and SaaS. Their approach leverage the WS-Agreement specification, a standard for SLA negotiation in web services, to define the SLA between the cloud provider and the cloud consumer [?]. Leveraging this representation, they propose a SLA-driven architecture with three main components: the **SLA Manager**, the **Orchestrator**, and the **Infrastructure Connector** [?]. The SLA Manager is the entry point for users and essentially builds and register agreements starting from the user requirements. The Orchestrator, having a complete view of all the available cloud backends, performs the critical task of scheduling and resource allocation, considering the SLA constraints. The Infrastructure Connector is the component that actually interacts with the cloud providers, performing the actual resource allocation and deallocation (provisioning and deprovisioning) [?]. An interesting feature of the Infrastructure Connector is the “configuration step” in which arbitrary actions can be performed such as the **injection of a monitoring agent on a Virtual Machine**.

In the context of serverless computing environments, the work of 2021 by Mampage et al. [?] proposes a deadline-aware dynamic resource management system. The focus of the research is on both the provider and the consumer perspective, proposing a placement policy and “dynamic resource management policy” that aims to minimize the cost of the provider while meeting the requirement of the consumer (i.e., the **deadline**) [?].

2.5 GreenOps landscape

GreenOps is the abbreviation for “**Green Operations**” and is the term used to describe an operational model that aims to integrate sustainable practices into an organization’s digital operations, with a particular focus on cloud computing and data centers. The interest in GreenOps has been growing in the last years due to the increasing awareness of the environmental impact of data centers and cloud infrastructures. As a matter of fact according to various sources, the carbon footprint of data centers is [...] The rising interest in the GreenOps ecosystem is also influenced by political and regulatory factors as the European Union’s ambitious goals for the reduction of greenhouse gas emissions. In particular, the *Climate Neutral Data Centre Pact in Europe*, is an EU initiative which aims for making climate-neutral data centers by 2030 [?]. Being the work of this thesis focused on the cloud ecosystem, we deem useful to cite the Technical Advisory Group (TAG) on Environmental Sustainability which is focused on the cloud-native sustainability landscape. The TAG is part of the Cloud Native Computing Foundation (CNCF) and aims to provide guidance, standards, and best practices for the industry [?].

2.5.1 Green Software foundation

Another pivotal actor in the GreenOps landscape is the Green Software Foundation. The Green Software Foundation is a non-profit organization, part of the Linux Foundation, that aims to promote the development of green software and to provide a set of standards, tooling and best practices for the industry [?]. It is considered useful to provide a quick summary of the foundation’s major projects that are relevant to the context of this work.

Software Carbon Intensity Specification

Software Carbon Intensity Specification (SCI) is a specification that aims to provide a standard way to measure the carbon intensity of software systems. SCI is defined as a rate: the amount of carbon emissions per one unit of R where R is a functional unit for the software system (e.g., API call, new user, DB query, etc.) [?].

The SCI can be defined as follows:

$$SCI = C \times R = (O + M) \times R = ((E \times I) + M) \times R$$

where:

- SCI is the Software Carbon Intensity
- C is the carbon emissions
- R is the functional unit
- O is the operational emissions
- M is the embodied emissions
- E is the energy consumption
- I is the carbon intensity

Real Time Cloud

The Real Time Cloud project aims to provide a standard for real-time carbon emissions data reporting for cloud providers. The goal is to provide real-time access to information about cloud region, PUE, WUE, carbon-free energy from the grid and from Cloud provider individual initiatives. The concept is similar to FOCUS specification for FinOps but is at a very early stage and is not yet adopted by cloud providers as of 2025 [?].

Impact Framework

The Impact Framework is a flexible framework for measuring and reporting the environmental impact of software systems [?]. The core of the framework is represented by a **Manifest file** which is a YAML file that is used both for describing **calculation pipelines** and for storing the results of the calculations. Pre-built plugins are available for the most common calculations (e.g., carbon intensity, energy consumption, etc.) but custom plugins can be developed as well [?]. A potential integration of this framework with our system is described in section 3.8.4.

2.5.2 Computational Sustainability by Public cloud providers

what are they already doing improving energy efficiency by reducing their PUE (Power Usage Effectiveness (PUE) is the ratio of the total energy used by a data center to the energy used for computing.)

For the purpose of this work, we assume that a cloud data center will likely rely on the same energy sources that characterize a specific geographical region (grid). For example, if data from Electricity Maps tell us that Finland is producing energy with low carbon emissions then we assume that the data centers in that area will likely be powered with energy from low carbon sources. However, some cloud providers may have better access to renewable energy sources in certain regions due to their individual initiatives e.g. wind farms that feed directly into their data centers.

what microsoft is already doing with alternative energy sources apart from grid

Carbon aware computing at Google <https://www.performance2021.deib.polimi.it/www.performance2021.deib.polimi.it/content/uploads/2021/11/Carbon-Aware-Computing-Year-2021> Description: It states that Carbon-aware computing is: exploiting flexibility in when and where and how computing is done to reduce carbon emissions. Some examples of flexible workloads are: video processing, training large-scale machine learning models, simulation pipelines. The components they recognized to be necessary are: accurate carbon intensity data (Tomorrow), Scalable infrastructures (Cloud), Virtualizations and migration mechanisms (VMs), Well identified flexible load, data-driven methodology. At Google: the total amount of work that needs to get done per day is quite predictable.

<https://blog.google/inside-google/infrastructure/data-centers-work-harder-sun-shines-wind-blows/>

Google CFE%: “This is the average percentage of carbon free energy consumed in a particular location on an hourly basis, while taking into account the investments we have made in carbon-free energy in that location. This means that in addition to the carbon free energy that’s already supplied by the grid, we have added carbon-free energy generation in that location”.

usually big companies leverage financial instruments like PPAs to buy renewable energy therefore their emissions are not zero but they are offset by the renewable energy they buy.

2.6 Carbon-aware systems for resource management

Having provided an overview of the existing works in the field of multi-cloud resource management and having introduced the GreenOps landscape, we now focus on the state of the art in the field of carbon-aware resource management. In particular, implementation choices and design patterns that can be leveraged for the development of a carbon-aware resource management system will be discussed.

The work of 2023 by Sukprasert et al. named “*Spatiotemporal Carbon-aware Scheduling in the Cloud: Limits and Benefits*” [?] is a comprehensive analysis on the limits and benefits of the employment of geographical shifting and time shifting for cloud workloads. The authors highlight the fact that different workloads have different characteristics and therefore different degrees of flexibility. Those include, for instance: **execution deadlines**, **data protection laws**, and **latency requirements**. Therefore, carbon savings are constrained by a complex set of factors that need to be taken into account when designing a carbon-aware system [?]. The following sections will present systems that have been developed by various research groups to tackle the problem of carbon-aware resource management in cloud computing environments.

2.6.1 CASPER

CASPER (Carbon-Aware Scheduling and Provisioning for Distributed Web Services) is a carbon-aware scheduling and provisioning system whose primary purpose is to minimize the carbon footprint of distributed web services [?]. The system is defined as a multi-objective optimization problem that considers two factors: the **variable carbon intensity** and the **latency constraints** of the network [?]. By evaluating the framework in real-world scenarios, the authors demonstrate that CASPER achieves significant reductions in carbon emissions (up to 70%) while meeting application **Service Level Objectives (SLOs)**, highlighting its potential for practical implementation in large-scale distributed systems [?]. The authors of CASPER highlight the importance of considering the workload characteristics such as memory state, **latency** and **regulatory constraints such as GDPR** [?]. The system is not adopting time-shifting since it is dealing with web services that are by their nature non-stopping workloads. The architecture is tied to scheduling K8s resources inside K8s clusters and does not consider external resource management.

2.6.2 CASPIAN

A research on carbon-aware scheduling in Kubernetes environments was conducted by the authors of CASPIAN (A Carbon-aware Workload Scheduler in Multi-Cluster Kubernetes Environments) [?]. The proposed system leverages Multi Cluster App Dispatcher (MCAD), a multi-cluster management platform, to provision workloads over distributed Kubernetes clusters [?]. Caspian is a scheduling and placement controller which lives in a master cluster and interacts with the MCAD to provision workloads across multiple geographical distributed clusters [?]. In particular, two main components are described: the **carbon tracker** and the **scheduler**. The carbon tracker is responsible for the periodic collection of carbon intensity data along with worker cluster locations. The scheduler, taking into account cluster information (carbon intensity, power efficiency, etc.), and workload requirements (e.g, **run time**, **deadline**, etc.), is responsible for the scheduling (time) and placement (geographical) of the workloads [?].

2.6.3 CarbonScaler

The work by Hanafy et al. “*CarbonScaler: Leveraging Cloud Workload Elasticity for Optimizing Carbon-Efficiency*” proposes a system that leverages the elasticity of batch cloud workloads to optimize carbon efficiency [?]. The targeted workloads are for instance **MPI jobs** and **Machine Learning training jobs**. The system is entirely Kubernetes-based and it is composed of three main components: the **Carbon Profiler**, the **Carbon AutoScaler**, and the **Carbon Advisor** [?]. The Carbon Profiler main duty is to estimate energy usage of jobs. The Carbon AutoScaler component is the core of the system and is a Kubernetes controller that leverage the **Kubeflow training operator**. Kubeflow is effectively used for the resource management of batch jobs. Finally, the Carbon Advisor simulates jobs execution in different configuration and allow the carbon reduction estimation [?].

2.6.4 A Low Carbon Kubernetes Scheduler

The work is focused on the extension of the Kubernetes scheduler (“kube-scheduler”) for the ranking and filtering of the “greenest” region for the deployment of entire Kubernetes clusters [?]. The system is tailored and tested on Azure but can be extended to other cloud providers. Provisioning operations of the clusters are done by the system leveraging Kubernetes API or IaaS management APIs [?]. Therefore Kubernetes is leveraged as a control plane to provision other Kubernetes clusters. An interesting feature proposed is the use of local air temperature and solar irradiance as tiebreaker for two datacenters with a similar carbon intense grid. The claim is that the solar irradiance has a bigger spread than the carbon intensity across global regions and that the local air temperature surrounding a datacentre affects the amount of energy needed for cooling [?].

2.6.5 Microsoft’s Carbon-Aware Kubernetes strategy

Citing to the previous work, Microsoft proposes a simple carbon-aware strategy for Kubernetes [?], integrating carbon intensity data into the scheduling process of Kubernetes Pods (rather than entire clusters). The Kubernetes scheduler, which allows custom rules for assigning Nodes to Pods, can incorporate carbon metrics like the Marginal Operating Emissions Rate (MOER) as factors in placement decisions. A weighted distribution can be created by normalizing the MOER values across the available Nodes. These weightings are encoded in a YAML file and applied as a priority for the Scheduler, which out-of-the-box supports these kind of custom rules to extend the scheduling process. The claim is that by combining three elements (i.e., Kubernetes scheduler, carbon intensity data, and a weighting algorithm) any Kubernetes instance can be made carbon-aware, at least in a simplified way [?].

2.7 Multi-cloud resource management - major takeaways

Many of the concepts described in this chapter are leveraged in the design and implementation of our system. We can briefly list some of the major recurrent elements that are generally needed in order to build a system for multi-cloud resource management:

1. representation of a generic workload (with its characteristics and constraints)
2. cloud service broker / orchestrator
3. modules to directly interact with the cloud providers

In the system proposed by this thesis, the elements listed above are represented by **cloud-native components** in the Kubernetes ecosystem as described in the following chapters.

3 System design and implementation

This chapter presents the design and implementation of our system, focusing on the integration of the various components and the overall architecture. A general description of some of the key components is provided to better explain their role in the system and the reasons for their inclusion. The system is designed to be modular, scalable, and extensible, enabling the integration of additional components as needed. It must be said the the following is a description of a first iteration of the system.

3.1 Assumptions

In this section we present the assumptions made during the design and implementation of the system.

3.1.1 Workload definition

In this work, workloads has been modeled as **Virtual Machines (VMs)**, representing the **primary use case** considered during the system's initial design phase. It is possible to define as “interruptible workloads”, those workloads that can be stopped and restarted without losing the work done. For this first iteration of the system only “**non-interruptible workloads**” are considered. This choice was driven by the fact the Virtual Machines are both a common and widely used cloud and one of the simplest cloud resources to provision and therefore ideal for the first iteration of the system. Having said that, for the purpose of this work, a formalization can be done.

A VM is defined as a tuple:

$$VM = (MinCPU, MinRAM, D, DL, ML)$$

where:

- $MinCPU$ is the minimum number of virtual CPUs required.
- $MinRAM$ is the minimum RAM required (in GB).
- D is the duration for which the VM must run to complete its processing task P (in hours).
- DL is the deadline timestamp by which the VM must complete execution.
- ML is the maximum allowed latency in milliseconds. If latency is not a constraint, then a high value can be set (e.g., $ML = 2000$).

This VM can be scheduled on any Public Cloud Provider since we are interested in a multi-cloud system where the cloud can be effectively seen as a commodity. Alternatively, a subset of **eligible Public Cloud Providers** can be set at runtime by the user. We will refer to **general-purpose VMs** and not specialized ones like GPU instances or high-performance computing instances.

As an example, we can define a VM with the following specifications:

$$VM_{example} = (4, 4, 2, "2025-03-20T23:59:59Z", 100)$$

where:

- $MinCPU = 4$ vCPUs
- $MinRAM = 4$ GB
- $D = 1$ hour
- $DL = "2025-03-20T23:59:59Z"$ (i.e., the processing task P inside the VM must complete before this timestamp)

The system is designed to be cloud-agnostic, however for the purpose of this work, the system is currently configured to support three major cloud providers: **Microsoft Azure**, **Google Cloud Platform (GCP)**, and **Amazon Web Services (AWS)**.

3.1.2 System limitations

A limitation of our general approach is that only resources supported by the cloud provider’s Kubernetes operator can be provisioned in a seamless way. Not all cloud resources available in a provider’s portfolio are guaranteed to have corresponding Kubernetes Custom Resources (CRs). This introduces certain constraints:

- Limited resource availability: if a specific resource type (e.g., a GPU-accelerated instance or a database service) is not supported by the cloud provider Operator, it cannot be provisioned using the current system.
- Dependence on Operator updates: cloud providers may extend or modify the set of resources supported by their Kubernetes operators over time.
- Vendor-specific implementations: for the same class of resources (e.g., virtual machines), the structure and fields of the CRs may vary a lot between cloud providers.

Despite these constraints, the system architecture remains highly adaptable, and future enhancements could incorporate additional or alternative provisioning mechanisms. An example of an alternative implementation could be the **direct API interactions** with cloud providers to bypass operator limitations. As a matter of fact, usually cloud provider Operators are leveraging these APIs under the hood to interact with the cloud provider’s services. Another approach could involve the development of custom operators or controllers to manage specific resource types not supported by existing operators. For instance, Krateo PlatformOps provides the so-called “oasgen-provider” that aims to fill the gaps of missing or incomplete Kubernetes operators. This module is a K8s controller that is able to generate Kubernetes CRDs and the related controllers from OpenAPI specifications [?].

3.2 System Architecture

The following table provides an overview of the main components of the system and their respective functions.

Component	Function
Krateo PlatformOps	Provides an abstraction layer for infrastructure orchestration, enabling declarative resource management and integration with cloud providers with templates.
Cloud Providers Kubernetes Operators	Manages the provisioning and reconciliation of cloud resources within Kubernetes, ensuring the actual state matches the desired state.
Kubernetes Mutating Webhook Configuration	Intercepts and modifies API requests before they are persisted, allowing dynamic configuration adjustments with policy enforcement.
OPA Server	Evaluates policy decisions based on defined constraints and input data from Kubernetes API requests through the webhook configuration.
OPA Policies and Data	Define the rules and contextual information used by OPA to make policy decisions, namely scheduling information
GreenOps Scheduler	Determines the optimal scheduling region and scheduling time for VMs, acting as an external data source for OPA policies.
MLflow	Allows the tracking, logging, versioning and storing of machine learning experiments for reproducibility and model lifecycle management.
KServe	Provides scalable and Kubernetes-native model serving capabilities, enabling deployment of machine learning models for inference.

Table 3.1: Main components of the system and their respective functions.

All the components listed in the above table must be deployed inside a Kubernetes cluster. The only exception are the OPA Policies and data which lies outside the cluster as described in section 3.6.6.

3.3 Krateo PlatformOps integration

Krateo PlatformOps is utilized in this system for **multi-cloud resource management**, allowing for the declarative orchestration of cloud resources across different cloud providers leveraging Kubernetes as a control plane. This section highlights the differences between two approaches that were considered for resource synchronization:

- The **Custom Kubernetes “Synchronization Operator”** approach (initially considered)
- The **Krateo PlatformOps** approach (adopted in the final system design)

It is deemed interesting to describe both approaches in order to identify the several **trade-offs** between implementing a custom synchronization operator and leveraging a template-based abstraction for cloud resource provisioning.

3.3.1 Resource management: the Custom Kubernetes “Synchronization Operator” approach

When dealing with multi-cloud workloads with Kubernetes as a control plane, a **synchronization and mapping mechanism** (broker) is required to bridge the gap between:

- **Generic Kubernetes Custom Resources**, which represent generic provider-agnostic workloads.
- **Cloud provider-specific Custom Resources**, which correspond to the actual cloud resources provisioned through the respective Kubernetes operators provided by Azure, AWS or GCP (in our case).

A custom Kubernetes Operator would be responsible for the mapping and synchronization of the above resource types. This approach is based on the principle of **Continuous Reconciliation**, where the operator continuously monitors and adjusts the system to maintain consistency between the desired and actual states. Candidate solutions for the development of a Kubernetes Operator includes: Operator SDK, Kubebuilder, or writing the operator from scratch using the Kubernetes client libraries. For the purpose of this work, Kubebuilder was used to develop a Proof of Concept (PoC) for the custom synchronization operator.

Responsibilities of the Custom Operator

In our specific case, the operator should continuously watch the generic CRs in the Kubernetes cluster to check if critical scheduling fields have been set:

- **schedulingRegion**: Defines where the workload should be placed.
- **schedulingTime**: Specifies when the workload should be deployed.

These fields, if set, indicate a geographical placement and timing for the workload have been determined by the GreenOps Scheduler. If these fields are not yet present, the operator must wait for scheduling decisions before proceeding. Therefore, inside its Reconcile() loop, the operator should:

1. Continuously check if scheduling fields (schedulingRegion, schedulingTime) are set.
2. Trigger the creation of the provider-specific resource when the schedulingTime is approaching.
3. Track the provisioning status by marking the generic CR with a field indicating that the cloud-specific resource has been created.

Post-Creation Considerations

Once the cloud provider-specific resource is created, two main questions arise:

- What happens if the provider-specific CR is modified manually?
- What happens if the VM configuration is modified directly on the cloud provider (outside Kubernetes)?

Related to the first question, an example scenario could be: changing the VM instance type (VM size) inside the Kubernetes cluster. In this case, the operator needs to decide whether to revert unauthorized changes or allow them and update the generic CR accordingly. For the second question, an example could be: changing the VM size directly on the cloud provider’s console. In this case, the operator should detect the drift and update the generic CR to reflect the external changes.

Resource linking

A mechanism must be in place to link the generic CR to the cloud provider-specific CR. Possible approaches include:

- UUID-based linking: A universally unique identifier ensuring each resource is mapped correctly.
- Kubernetes Object Metadata (ObjectMetadata.Name & ObjectMetadata.Namespace): This approach may be preferable within a single Kubernetes cluster, avoiding the need for an external ID system.

Termination Logic

The operator must handle the deletion of cloud resources correctly in a variety of scenarios, including:

- When the provider-specific CR is deleted from Kubernetes, the corresponding cloud resource is de-provisioned and the custom operator should ensure the deletion process is handled gracefully, avoiding orphaned generic CRs.
- If the provider-specific resource is deleted directly on the cloud provider (e.g., on the provider console), the operator should detect the change and update the generic CR accordingly.
- In the event of a generic CR deletion, the custom operator should ensure the provider specific resource is removed, triggering a deletion process on the cloud provider side (de-provisioning).

Managing cloud provider-specific fields

Each cloud provider has unique resource configurations and constraints that must be managed. Some differences are purely syntactic (e.g., AWS uses *instanceType*, whereas Azure uses *vmSize*). Others require additional provider-specific metadata (e.g., Azure requires a *resourceGroup* field which represent a logical container for resources in Azure). A custom synchronization operator must take into account and encode this logic explicitly, making it more complex to maintain especially when supporting several cloud providers.

Limitations of a Custom “Synchronization Operator”

Another major challenge with a custom synchronization operator is that some cloud providers do not support time scheduling metadata within their Custom Resources. In particular, no cloud provider operator among the ones we used for the system (AWS, Azure, GCP) currently provides a dedicated field for scheduling time. This means that the custom Kubernetes operator itself must handle a time scheduling logic, delaying CR creation until the scheduled time. If the operator, upon the creation of a generic CR, immediately creates the cloud-provider specific CR (without a “waiting logic”), the cloud provider Operator will trigger and provision the VM immediately, ignoring scheduling constraints. Due to these limitations and complexities, we explored and leveraged an **alternative template-driven approach** using Krateo PlatformOps, described in the next section.

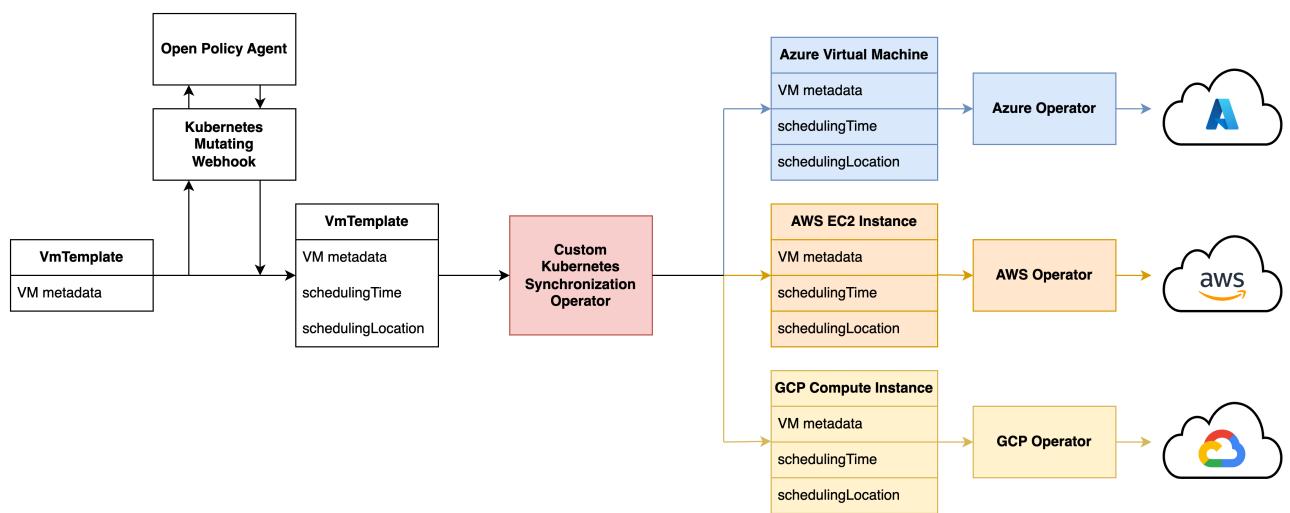


Figure 3.1: Multi-cloud resource management with Custom Kubernetes “Synchronization Operator” approach

3.3.2 Resource management: the Krateo PlatformOps approach

In our approach, we opted to replace a custom Kubernetes operator (“Synchronization Operator”), originally designed to handle the **mapping** from generic to cloud-specific resources, with **Krateo Core Provider**. This decision was motivated by the need for **greater flexibility and maintainability**.

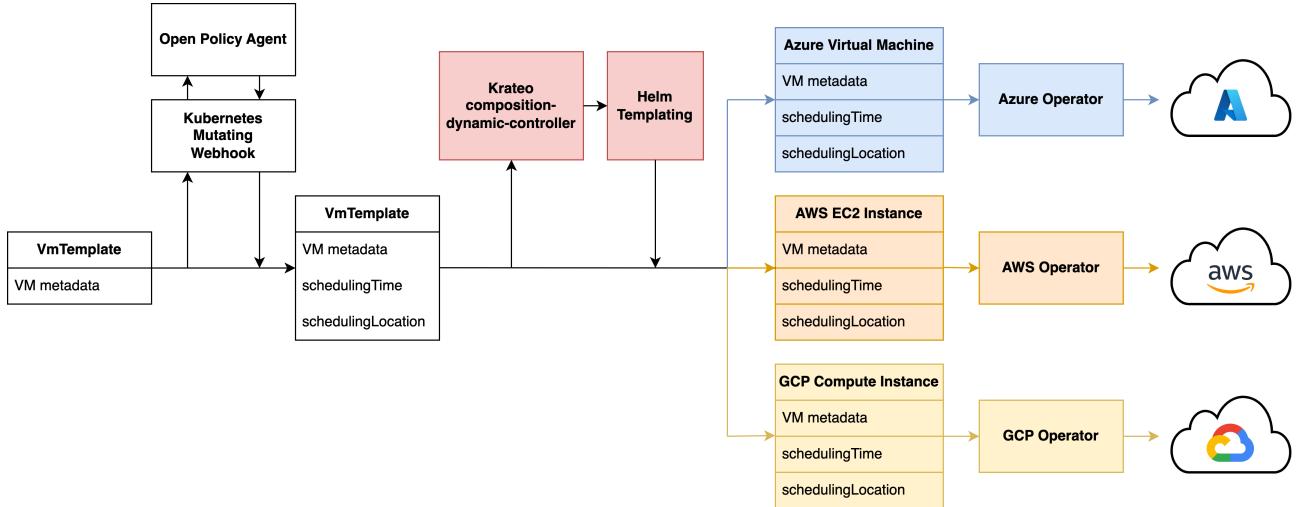


Figure 3.2: Multi-cloud resource management with Krateo PlatformOps approach

in defining multi-cloud infrastructure components. As a matter of fact, a custom Kubernetes operator was originally designed to handle only virtual machines (VMs). Mappings and extensions to support additional cloud resources would have required significant code changes and maintenance overhead for each additional resource type added. Therefore instead of embedding business logic directly within a custom Kubernetes operator, in the current system implementation, we leverage the capabilities of **Helm templating** to dynamically generate cloud-provider-specific resources. More precisely, another Krateo component, the **Krateo composition-dynamic-controller** is leveraging **Helm Template Engine** under the hood to generate Kubernetes resources starting from Helm templates. By adopting an Helm-based resource generation, we can reduce the maintenance overhead and simplify the system architecture. Figure 3.2 illustrates, in a high-level manner, the revised system architecture for resource management with Krateo PlatformOps.

This approach, further described in this section, offers several advantages:

- Simplified resource management: Helm enables a standardized way define resources without maintaining complex operator logic.
- Greater extensibility: By externalizing the logic from the Custom Operator (effectively removing it), future modifications and integrations on the system with additional cloud providers become easier.
- Reduced maintenance overhead: Custom operators typically require constant updates and refinements, especially if they are responsible for complex business logic.

Generic VM resource definition

In order to define a generic VM resource, we leverage Krateo Core Provider to define a **Composition-Definition** resource that specifies the structure of the VM resource. Effectively a CompositionDefinition is a Kubernetes Custom Resource that defines the structure of a Composition Custom Resource (instance) which is an Helm chart comprised of Helm templates and values. As a matter of fact, a CompositionDefinition is pointing to a versioned Helm chart that is stored in a Helm repository (e.g. on a Helm repository server).

```

1 apiVersion: core.krateo.io/v1alpha1
2 kind: CompositionDefinition
3 metadata:
4   name: vmtemplate
5 spec:
6   chart:
7     repo: vm-template
8     url: https://leonardovicentini.com/helm-charts/charts
9     version: 1.2.0

```

Listing 3.1: CompositionDefinition

The Helm chart referenced in the CompositionDefinition contains the Helm templates and values needed to define the resource named *VmTemplate*. Listing 3.2 shows an example of the VmTemplate values.yaml file, which defines the fields of the VM resource while listing ?? shows the corresponding JSON schema for the values.yaml file.

```
1 # @param {string} vmName Name of the VM
2 vmName: test-vm
3
4 # @param {integer} cpu Number of CPU cores
5 cpu: 1
6
7 # @param {integer} memory Number of GB of RAM
8 memory: 2
9
10 # @param {string} [schedulingTime] Scheduling Time for the VM
11 schedulingTime: 2025-05-05T00:00:00Z
12
13 # @param {string} [schedulingLocation] Scheduling Location for the VM
14 schedulingLocation: italynorth
15
16 # @param {string} duration Duration of the Workload
17 duration: 3h
18
19 # @param {string} deadline Deadline of the Workload
20 deadline: 2025-12-31T10:00:00Z
21
22 # @param {integer} maxLatency Maximum Latency of the Workload
23 maxLatency: 100
```

Listing 3.2: values.yaml

```

1  {
2      "type": "object",
3      "$schema": "http://json-schema.org/draft-07/schema",
4      "required": [
5          "vmName",
6          "cpu",
7          "memory",
8          "duration",
9          "deadline",
10         "maxLatency"
11     ],
12     "properties": {
13         "vmName": {
14             "type": [
15                 "string"
16             ],
17             "description": "Name of the VM",
18             "default": "test-vm"
19         },
20         "cpu": {
21             "type": [
22                 "integer"
23             ],
24             "description": "Number of CPU cores",
25             "default": "1"
26         },
27         "memory": {
28             "type": [
29                 "integer"
30             ],
31             "description": "Number of GB of RAM",
32             "default": "2"
33         },
34         "schedulingTime": {
35             "type": [
36                 "string"
37             ],
38             "description": "Scheduling Time for the VM",
39             "default": "2025-05-05T00:00:00Z"
40         },
41         "schedulingLocation": {
42             "type": [
43                 "string"
44             ],
45             "description": "Scheduling Location for the VM",
46             "default": "italynorth"
47         },
48         "duration": {
49             "type": [
50                 "string"
51             ],
52             "description": "Duration of the Workload",
53             "default": "3h"
54         },
55         "deadline": {
56             "type": [
57                 "string"
58             ],
59             "description": "Deadline of the Workload",
60             "default": "2025-12-31T10:00:00Z"

```

```

61 },
62   "maxLatency": {
63     "type": [
64       "integer"
65     ],
66     "description": "Maximum Latency of the Workload",
67     "default": "100"
68   }
69 }
70 }
```

Listing 3.3: values.schema.json

Generic VM to Cloud Provider Specific VM mapping

The Krateo composition-dynamic-controller (Krateo cdc) is responsible for creating the Composition Custom Resource (instance) by templating the Helm chart referenced in the CompositionDefinition. For the purpose of the system, we leveraged Helm templating to dynamically generate cloud-provider-specific resources from a generic VM resource. As a matter of fact, the Helm chart contains all the **three cloud provider specific set of templates** necessary for a VM provisioning but **only one set of templates is used at a time**. The Krateo cdc will generate the cloud-provider-specific VM resources based on the cloud provider specified in the generic VM resource. Therefore, for each generic VM resource, the Krateo cdc will use **only one set of templates** to generate the cloud provider specific VM resource. This “brokering mechanism” (routing through Helm templates) is implemented in the Helm templates using “guards” as reported in listing 3.4. The manifest will be created only if the cloud provider specified in the generic VM resource is the same as the cloud provider specified in the generic VM resource. This flexibility allows the system to be cloud-agnostic and to support multiple cloud providers as soon as the cloud provider specific templates are available in the Helm chart. The directory structure of the Helm chart is the following:

```

chart/
  [several utilities files for mappings (omitted)]
  templates/
    aws/
      instance.yaml
      subnet.yaml
      vpc.yaml
    azure/
      networkinterface.yaml
      virtualmachine.yaml
      virtualnetwork.yaml
      virtualnetworksubnet.yaml
    gcp/
      computeinstance.yaml
      computenetwork.yaml
      computesubnetwork.yaml
      helpers.tpl
    Chart.yaml
    values.schema.json
    values.yaml
```

```

1 {{ if hasKey .Values "provider" }}
2 {{ $provider := .Values.provider }}
3 {{ if eq $provider "azure" }}
4 ...
```

```

5 apiVersion: compute.azure.com/v1api20220301
6 kind: VirtualMachine
7 metadata:
8   name: {{ .Values.vmName }}
9 ...

```

Listing 3.4: Helm Template guards example

VM size selection is a crucial step in the VM provisioning process. Helm allows the definition of **helper functions** which resides in the `_helpers.tpl` file. In our case, we defined a helper function called `findBestVmSize()` that takes as input the CPU and RAM requirements of the VM and returns the best provider-specific VM size available. For instance: a generic requested VM specification could be (4 vCPU, 8 GiB of RAM) and this would be mapped to the “*Azure Standard_A4_v2*” VM size. A prerequisite for that is to have a mapping between the tuple (CPU, RAM) and the VM sizes (a string) for each cloud provider. This mappings are were built using the cloud provider documentation and are stored in the Helm chart as a set of utilities files. Potentially, an automatic way to fetch the available VM sizes for each cloud provider could be implemented in the future. For our first iteration of the system we used a small subset of all the available instance types (i.e. 5 instance types for each cloud provider).

Scheduling time waiting logic

Helm template engine is also leveraged to handle the **scheduling time waiting logic**. As previously described in section 3.3.1, this logic is crucial for the system due to the fact that cloud provider operators do not support scheduling time metadata in their CRs and therefore as soon as a CR is created the cloud provider operator will provision the resource immediately. It must be remembered that is not a trivial task to implement this logic in a Kubernetes operator. In this case, we leverage a set of “**guards**” to effectively block manifest creation until the scheduling time is reached. As a result, only when the scheduling time is reached the manifest is created and applied by Helm and finally the resource is provisioned by the cloud provider operator. As a matter of fact, Krateo cdc will periodically perform an “helm upgrade” and if the scheduling time is reached the cloud provider specific resources will be created. As a consequence, the specific cloud provider operator will be triggered by the creation of the cloud provider specific resources and will provision the VM. Listing 3.5 shows an example of how guards are used to handle the scheduling time waiting logic and other scheduling constraints. We also make sure, for instance, that the generic CR contains a “provider” field that specifies the cloud provider where the resource should be provisioned.

```

1 {{ if hasKey .Values "provider" }}
2 {{ $provider := .Values.provider }}
3 {{ if eq $provider "azure" }}
4
5 {{ if hasKey .Values "schedulingTime" }}
6 {{ $schedulingTime := .Values.schedulingTime | toDate "2006-01-02T15:04:05Z" }}
7 {{ $now := now }}
8 {{ if $now.After $schedulingTime }}  

9 {{ if $now.After $schedulingTime }}
10 ...
11 apiVersion: compute.azure.com/v1api20220301
12 kind: VirtualMachine
13 metadata:
14   name: {{ .Values.vmName }}
15 ...

```

Listing 3.5: Scheduling time

In order to take into account the time needed for the cloud provider operator to provision the resource, additional logic could be added in the template to anticipate the scheduling time (e.g, using a `provisioningTimeOffset`).

3.4 Multi-Cloud Integration through Kubernetes Operators

The integration of operators from different cloud providers has enabled the development of an effective **multi-cloud system**, allowing seamless orchestration and provisioning of cloud resources across various cloud platforms. More precisely, the system leverages Kubernetes operators from **Microsoft Azure**, **Google Cloud Platform (GCP)**, and **Amazon Web Services (AWS)**. Each Operator, when installed on a Kubernetes cluster, installs a **set of Custom Resource Definitions (CRDs)** that represent cloud resources specific to the cloud provider. These CRDs can be used to define cloud resources in a declarative manner, in the form of Kubernetes Custom Resources (CRs), allowing users to specify the desired state of the cloud resources they wish to manage. The other component installed by the Operator is the **controller**, a software module that watches for changes to the Custom Resources in the cluster and takes all the necessary actions to reconcile the actual state with the desired state. Under the hood the controller interacts with the cloud provider's API to provision, update, and delete cloud resources. Kubernetes operators work on the principle of **Continuous Reconciliation**, ensuring, in this case, that the desired state of the system, as defined by users, aligns with the actual state of provisioned cloud resources. In particular, Operators act as controllers that monitor (*watch*), adjust, and manage external cloud resources within a Kubernetes environment. Inside the Kubernetes cluster lie the **real-time representation of the provisioned cloud resources**, which are managed by the operators. It must be noted that different cloud providers adopt **different design choices** for their Kubernetes operators and more in general for their overall cloud infrastructure management. Therefore, for the creation of logically similar resources, like a virtual machine, the structure and the field of the resources can be different. These resources typically include:

- Compute resources (e.g., VM instances, virtual machine templates)
- Networking components (e.g., virtual networks, subnets, security groups)
- Storage allocations (e.g., persistent volumes, cloud disks)
- Access management (e.g., resource groups, roles, authentication credentials)

For the purpose of this work we defined a **baseline infrastructure** for each cloud provider taken into account in order to have a common ground for the system to work. This baseline infrastructure is composed by the minimum set of resources needed for a VM provisioning. Each public cloud provider has its complexities and nuances when it comes to managing cloud resources. In the following sections, we provide an overview of the minimum set of resources needed for VM provisioning on each cloud provider, as well as some of the specific configurations required for each resource. We deem useful to provide a table summarizing some of the key fields that are needed for VM provisioning. These fields are labeled with an ID which will be attached to the resources' listings in the following sections.

ID	VM field
1	Scheduling region
2	VM size
3	Operating System

Table 3.2: Key fields for VM provisioning

3.4.1 Azure Kubernetes Operator

Microsoft Azure provides a Kubernetes operator called **Azure Service Operator v2** (ASO). Currently, ASO supports more than 150 different Azure resources one of which is the **Azure Virtual Machine**. The minimum set of resources needed for VM provisioning on Azure through Azure service operator is:

- Virtual Network
- Virtual Network Subnet
- Network Interface
- Virtual Machine

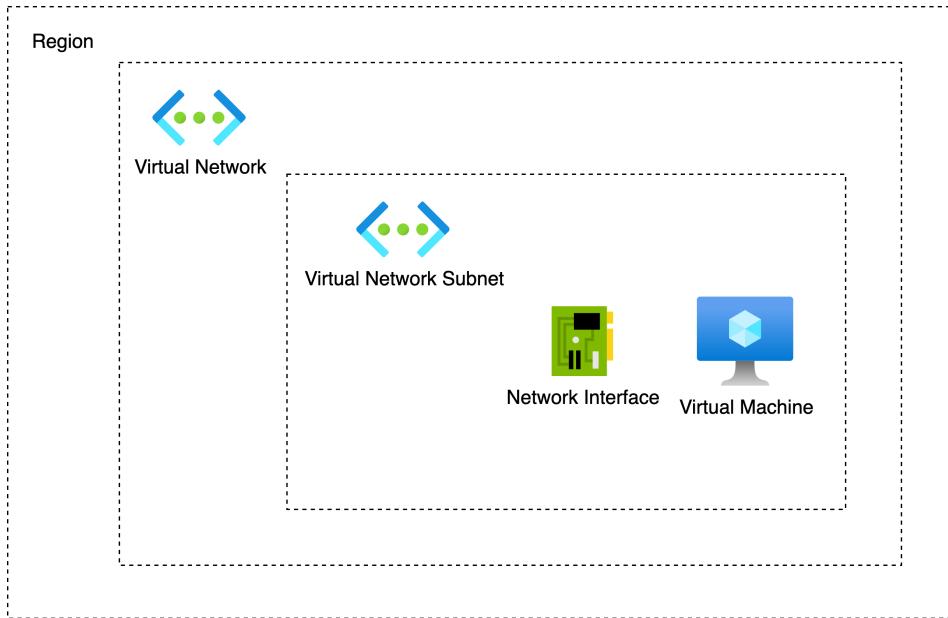


Figure 3.3: Minimum set of Azure resources for VM provisioning

An example of a Azure-specific custom resource is the one illustrated in listing 3.6.

```

1 apiVersion: compute.azure.com/v1api20220301
2 kind: VirtualMachine
3 metadata:
4   name: {{ .Values.vmName }}
5   namespace: {{ .Values.namespace | default "greenops" }}
6 spec:
7   hardwareProfile:
8     vmSize: {{ $vmSize }} #[2]
9     location: {{ .Values.schedulingLocation }} #[1]
10 networkProfile:
11   networkInterfaces:
12   - reference:
13     group: network.azure.com
14     kind: NetworkInterface
15     name: {{ .Values.vmName }}-{{ $av.networkInterface.suffix }}
16 osProfile: {{ $av.virtualMachine.osProfile | toYaml | nindent 4 }}
17 owner:
18   armId: {{ $av.resourceGroup.armId }}
19 storageProfile:
20   imageReference: #[3]
21     publisher: Canonical
22     offer: 0001-com-ubuntu-server-jammy
23     sku: 22_04-lts
24     version: latest

```

Listing 3.6: Azure Instance Custom Resource

3.4.2 GCP Operator

Google Cloud Platform provides a Kubernetes operator called **GCP Config Connector**. The name of the virtual machine resource in GCP is **ComputeInstance**. For what concerns the GCP Operator, the minimum set of resources needed for VM provisioning is:

- Compute Network
- Compute SubNetwork
- Compute Instance

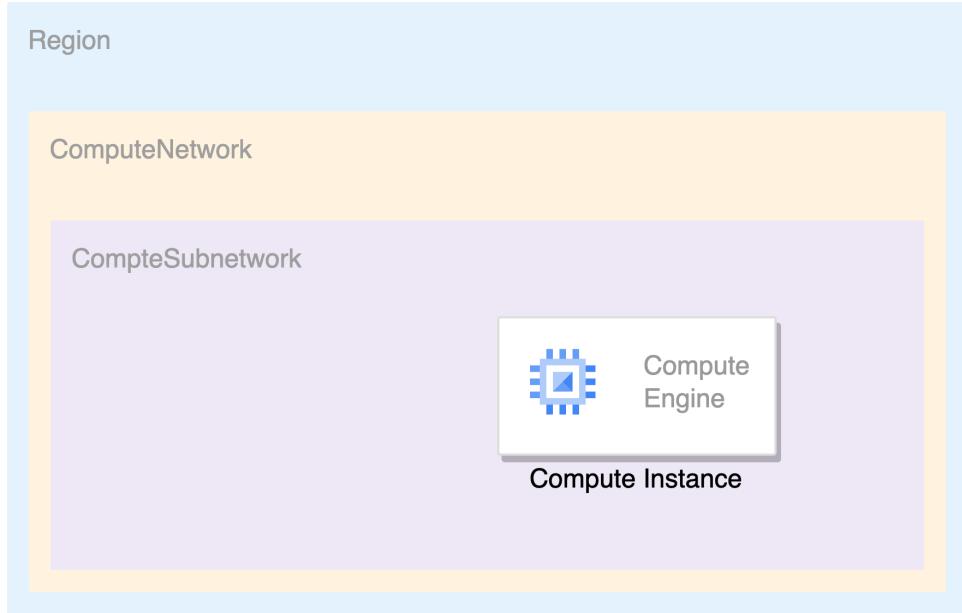


Figure 3.4: Minimum set of GCP resources for VM provisioning

```

1  apiVersion: compute.cnrm.cloud.google.com/v1beta1
2  kind: ComputeInstance
3  metadata:
4    name: {{ .Values.vmName }}
5    namespace: {{ .Values.namespace | default "greenops" }}
6  spec:
7    machineType: {{ $vmSize }} #[2]
8    machineType: {{ $vmSize }} [2]
9    zone: {{ $zone }} #[1]
10   zone: {{ $zone }} [1]
11   bootDisk:
12     initializeParams:
13       size: 24
14       type: pd-ssd
15     sourceImageRef:
16       external: debian-cloud/debian-11 #[3]
17       external: debian-cloud/debian-11
18   networkInterface:
19     - subnetworkRef:
20       name: {{ .Values.vmName }}-subnetwork
21     aliasIpRange:
22       - ipCidrRange: /24
23       subnetworkRangeName: cloudrange

```

Listing 3.7: GCP Intance Custom Resource

3.4.3 AWS Operator

AWS provides an entire collection of operators that are part of the AWS controllers for Kubernetes (ACK) project. Each controller is responsible for managing a specific AWS service, such as EC2, EBS, RDS, S3, and more. In the context of this research, the EC2 controller is used to manage virtual machine (EC2 Instances) provisioning on AWS. The minimum set of resources needed for VM provisioning is:

- VPC
- Subnet
- Instance

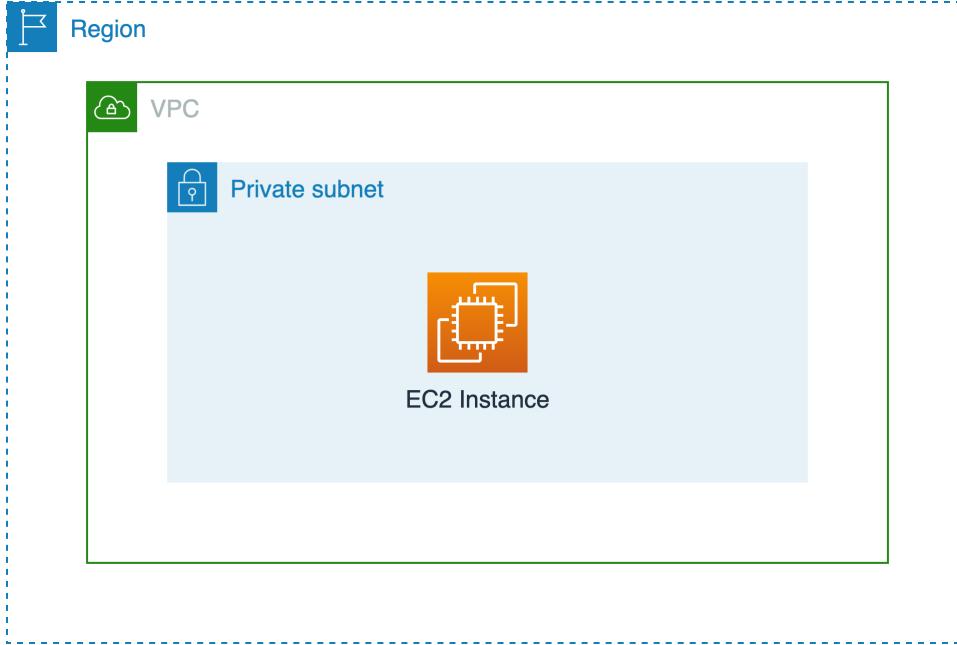


Figure 3.5: Minimum set of AWS resources for VM provisioning

As described at the beginning of this section, the implementation approach adopted in our system ensures compatibility with diverse cloud provider design choices. Cloud providers may impose different constraints and best practices when managing Kubernetes-native resources, and the system is designed to adapt to these variations seamlessly. One notable design choice observed with the AWS operator is the restriction on referencing some Kubernetes objects inside a Custom Resource (CR) manifest. This limitation means that developers cannot directly link a resource (e.g., a Virtual Machine) to another Kubernetes object (e.g., a Subnet) using built-in object references. To overcome this limitation, our system leverages **Helm's *lookup function***, which dynamically retrieves Kubernetes object details at runtime. This method allows us to fetch required parameters without directly referencing Kubernetes objects in the CR, ensuring compatibility with the AWS operator's design constraints. The following example demonstrates how the lookup function can be used to resolve subnet IDs dynamically and inject them into the CR manifest.

```

1 ...
2 apiVersion: ec2.services.k8s.aws/v1alpha1
3 kind: Instance
4 metadata:
5   name: {{ .Values.vmName }}
6   namespace: {{ .Values.namespace | default "greenops" }}
7 ...
8 spec:
9 ...
10   subnetID: {{ (lookup "ec2.services.k8s.aws/v1alpha1" "Subnet" (.Values.namespace | default "greenops")) (printf "%s-subnet" .Values.vmName)).status.subnetID }}
11 ...

```

Listing 3.8: Helm Lookup example: dynamically resolving SubnetIDs

The Helm lookup function can be used to look up resources in a running cluster and its synopsis is: “`lookup apiVersion, kind, namespace, name -j` resource or resource list” [?]. In the listing 3.8, the Helm lookup function retrieves the subnetID from a Subnet Custom Resource dynamically, based on the VM name and namespace. Then, the subnetID is injected into the Instance Custom Resource manifest, ensuring that the VM is provisioned in the correct subnet. An example by the same AWS Operator where instead a direct reference to a resource is allowed is the one illustrated in listing 3.9.

```

1 ...
2 apiVersion: ec2.services.k8s.aws/v1alpha1
3 kind: Subnet
4 metadata:
5   name: {{ .Values.vmName }}-subnet
6 ...
7 spec:
8   vpcRef:
9     from:
10    name: {{ .Values.vmName }}-vpc
11    namespace: {{ .Values.namespace | default "greenops" }}
12 ...

```

Listing 3.9: AWS Operator direct reference example

In the case of Listing 3.9, the Subnet Custom Resource manifest directly references in a convenient way the VPC Custom Resource using its name and namespace since the Operator is designed to support this type of relationship. As explained before, this is determined by Operator design choices but our system is able to handle both scenarios.

Provider specific configurations

When launching an Amazon EC2 instance, specifying an AMI is **mandatory**. An Amazon Machine Image (AMI) is a pre-configured image that provides the necessary software environment to set up and boot an Amazon EC2 instance [?]. In other words, AMIs serve as a blueprint for launching virtual machines (VMs) in AWS. The AMI must be compatible with the chosen EC2 instance type, ensuring that the selected image supports the required hardware and software configurations. The following attributes define an AMI:

- Region: AMIs are region-specific
- Operating System: Determines the base OS (e.g., Ubuntu, RHEL) installed on the AMI.
- Processor Architecture: e.g., x86, ARM
- Root Device Type: Specifies whether the AMI uses an EBS-backed volume (Elastic Block Store) or Instance Store for storage.
- Virtualization Type: Defines whether the AMI supports paravirtual (PV) or hardware virtual machine (HVM) instances.

The most important attribute for our system is the **Operating System** since it determines the software environment for the VM. For the purpose of this research, only **Ubuntu-based AMIs** have been considered for provisioning virtual machines. Official Ubuntu AMIs were collected from a dedicated Ubuntu repository. In order to select the most suitable AMI for a given VM, the system leverages a function that

to dynamically select the appropriate AMI ID based on the region and other parameters specified in the VmTemplate Kubernetes Custom Resource (CR).

```

1 apiVersion: ec2.services.k8s.aws/v1alpha1
2 kind: Instance
3 metadata:
4   name: {{ .Values.vmName }}
5   namespace: {{ .Values.namespace | default "greenops" }}
6   annotations:
7     services.k8s.aws/region: {{ .Values.schedulingLocation }} #[1]
8 spec:
9   imageID: {{ $imageID }} #[3]
10  instanceType: {{ $vmSize }} #[2]
11  subnetID: {{ (lookup "ec2.services.k8s.aws/v1alpha1" "Subnet" (.Values.namespace | default "greenops")) (printf "%s-subnet" .Values.vmName)).status.subnetID }}

```

Listing 3.10: AWS

3.5 Kubernetes Mutating Webhook Configuration

In the context of **Kubernetes Admission Control**, in addition to standard, compiled-in admission plugins, Kubernetes supports the use of additional admission plugins that are effectively extensions of the system and run as **webhooks** configured at runtime [?]. This means that the admission control logic can be extended dynamically without the need to recompile the Kubernetes API server or other Kubernetes components. Changes are applied at runtime to the running Kubernetes cluster, making the system more flexible and adaptable. Said plugins can be used to enforce custom policies and perform additional validation and mutation of Kubernetes objects before they are persisted in the cluster. We can classify webhooks in two categories: validating and mutating webhooks. Validating webhooks are used to validate the object and reject it if it does not meet the validation criteria. Mutating webhooks are used to modify the object before it is persisted in the cluster. We must highlight the difference of two entities: the webhook configuration and the actual webhook server which performs mutation or validation. The latter could be a custom server to apply custom mutation or validation logic or it could be a service ready to use out of the box like an Open Policy Agent server.

Figure 3.6 shows the Kubernetes Admission Control flow with the addition of mutating and validating webhooks. The flow can be summarized as follows:

1. The K8s API server receives an API request.
2. Authentication and authorization phase take place.
3. Mutating webhooks are called first. If any of them returns an error, the request is rejected.
4. The Object Schema Validation phase is executed.
5. Validating webhooks are called. If any of them returns an error, the request is rejected.
6. The object is persisted in the cluster.

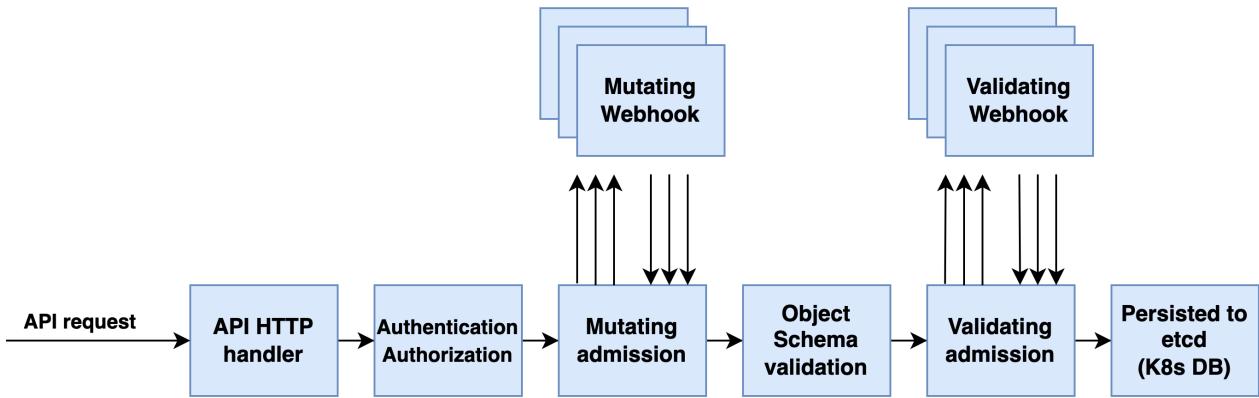


Figure 3.6: Kubernetes Admission Control

A webhook service can be either a deployment with the related service inside the cluster like seen in the image 3.7 or can be a service deployed outside the cluster as long as it is reachable by the K8s API server.

Figure 3.7 shows an example of a Kubernetes Mutating Webhook. The webhook server is deployed as a Kubernetes Deployment (with 1 Pod), with a corresponding Service to expose it within the cluster. The webhook server is responsible for receiving The AdmissionReview Request, applying custom logic, and returning an AdmissionReview Response to the Kubernetes API server describing the mutation to be applied to the resource. The custom logic in this simple example is to add a label "mutated: true" to the resource.

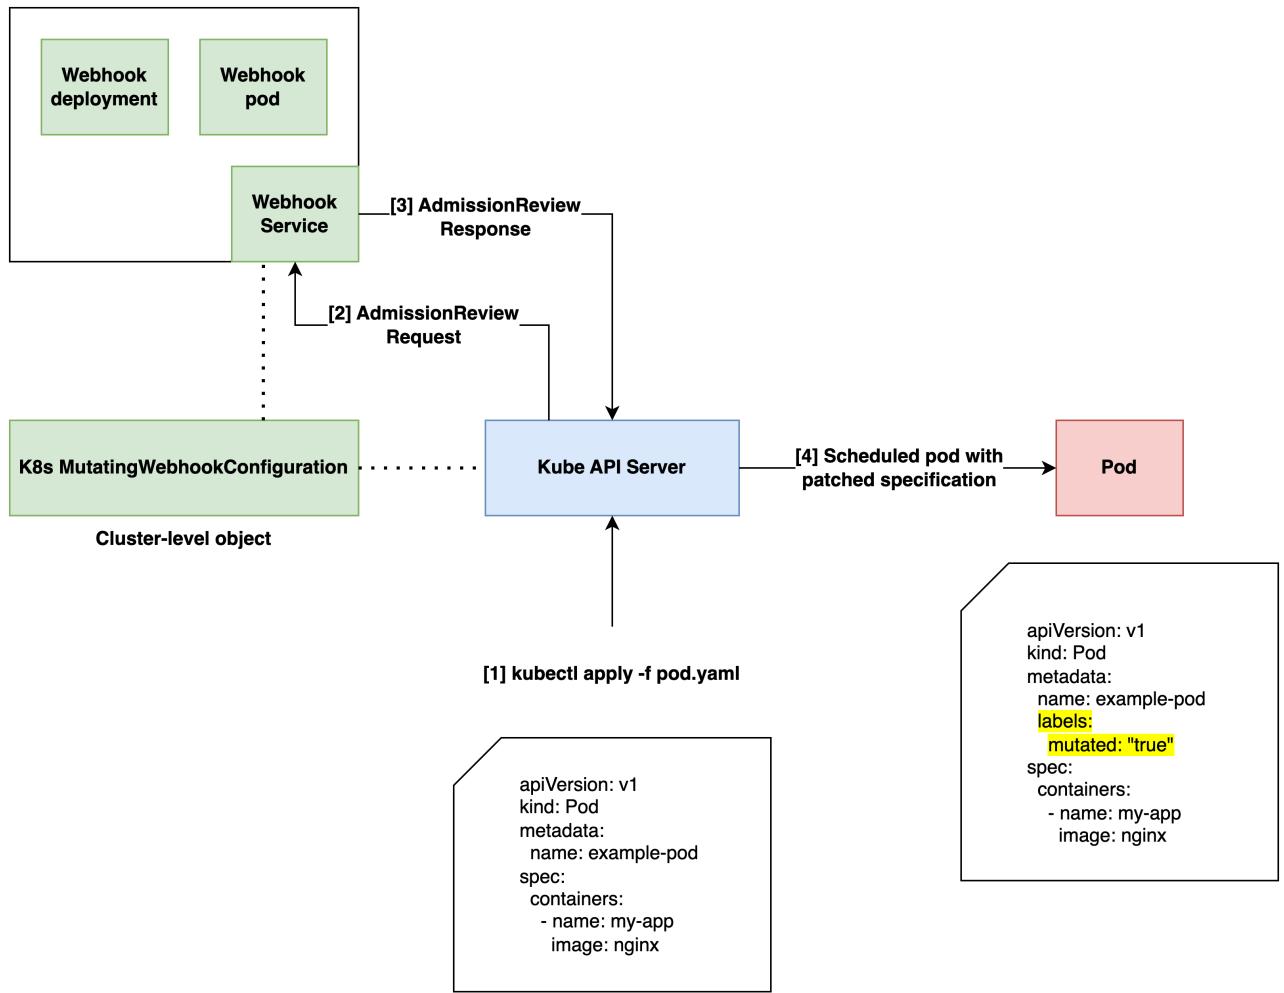


Figure 3.7: Kubernetes Mutating Webhook example [?]

In the context of our system, a **Mutating Webhook Configuration** is set to intercept CREATE and UPDATE Kubernetes API requests for VmTemplate Custom Resources. In this case, the Kubernetes Mutating Webhook server that is the target of the MutatingWebhookConfiguration is Open Policy Agent (OPA). OPA is used to assign scheduling decisions to VmTemplates based on policies. OPA sends back to the API server the mutation (patches) to be applied to the VmTemplate Custom Resource. This process is further described in the next section.

3.6 Open Policy Agent (OPA)

Open Policy Agent (OPA) is an open-source general-purpose **policy engine** that enables unified policy enforcement across several types of environments. OPA provides a declarative language called **Rego** enabling a paradigm known as “**Policy as Code**” [?].

Open Policy Agent can be integrated as a sidecar container, host-level daemon, or library to perform policy decisions for a plethora of use cases: microservices, Kubernetes admission control, CI/CD pipelines, API gateways and more [?].

In the context of our system, OPA and the Policy-as-Code paradigm are leveraged to define **policies for workload scheduling**: encoding the output of a **scheduling decision** coming from an external GreenOps Scheduler and ensuring compliance with **additional policies** related to latency requirements and legal constraints.

3.6.1 Policy as Code paradigm

According to AWS, Policy-as-Code (PaC) is a software automation approach which is similar to Infrastructure-as-Code (IaC) [?]. PaC helps assess company system configurations and validate compliance requirements through software automation [?]. The perceived value of this type of automation

in the software development lifecycle has grown significantly in modern enterprises. This large adoption is probably driven by the inherent consistency and reliability it provides, ensuring standardized enforcement of policies and reducing human error [?].

OPA's generic definition of policy is: “*A policy is a set of rules that governs the behavior of a software service*” [?]. OPA provides a high-level declarative language called **Rego** to define policies in a flexible manner. One of OPA’s key strengths is its **domain-agnostic design**, allowing it to enforce policies across various systems and environments. This makes it highly adaptable to different use cases, ranging from access control to infrastructure security. Some representative examples of policies that OPA can enforce include:

- Restricting which image registries can be used for deploying new Pods in a Kubernetes cluster.
- Controlling whether a specific user is permitted to perform delete operations on certain resources.
- Enforcing network security policies, such as blocking external access to sensitive services.
- Ensuring infrastructure compliance, for example, by verifying that new cloud resources to be provisioned follow predefined security configurations.
- Enforcing that new deployed servers must have the prefix “server-” in their name.

Another interesting example is the use of OPA for compliance automation for AWS infrastructure [?]. In this case, OPA is used as a step in a CI/CD pipeline to enforce compliance policies on AWS infrastructure represented with the Infrastructure as Code (IaC) tool CloudFormation [?].

Therefore, the use cases covered span from role-based access control to container image security and beyond.

Another important aspect of OPA is that it effectively **decouples** policy decision-making from policy enforcement, enabling organizations to implement consistent and scalable authorization across their systems [?]. In practice, this means that when a software module needs to make a policy decision, it queries OPA, supplying relevant data as input. In other words, policy decisions are **offloaded** to OPA rather than being hardcoded within individual services. This approach offers several key advantages:

- **Centralized policy management:** policies are defined in a single location, ensuring uniform enforcement across all services of an organization.
- **Improved maintainability:** updating policies does not require modifying, recompiling or redeploying application code, reducing complexity and deployment overhead.
- **Greater flexibility:** policies can be dynamically updated (e.g., with CI/CD approaches) based on evolving security and compliance requirements.
- **Scalability:** since OPA and application modules are not tightly coupled.

3.6.2 OPA architecture overview

As mentioned in the introduction to this section, one common approach to integrating OPA into a software system is by deploying it as a host-level daemon. The latter is essentially a lightweight server that processes policy queries via HTTP requests. This setup allows services to offload policy decision-making to OPA in a scalable and efficient manner since the two entities are not tightly coupled.

A standard OPA deployment consists of three main components:

- **OPA Server:** The core service that evaluates policy queries and returns decisions based on defined rules, contextual data and input data.
- **OPA Policies:** Rules written in the Rego language that define the logic to be enforced.
- **Data:** Optional contextual information, typically structured in JSON format, that policies use to make informed decisions along with input data.

To facilitate deployment and management, Rego policies and associated contextual data are packaged into **policy bundles**, as described in section 3.6.5. These bundles enable version-controlled, centralized policy distribution, ensuring consistency and maintainability across distributed environ-

ments.

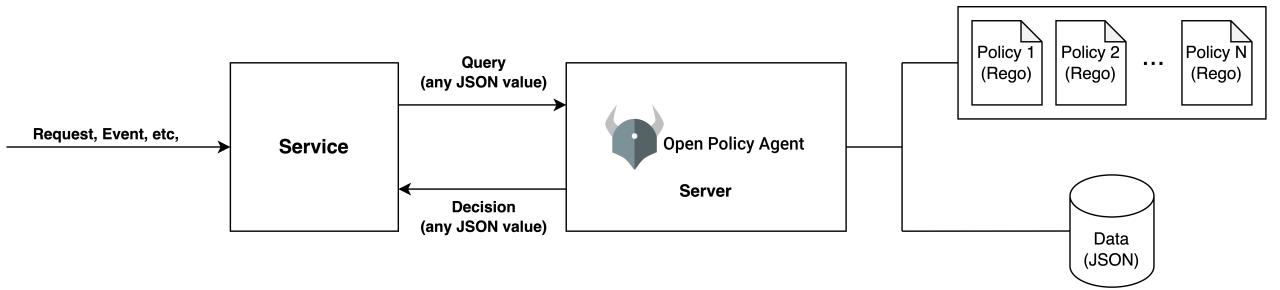


Figure 3.8: OPA architecture

OPA accepts arbitrary structured data as input. and Like query inputs, your policies can generate arbitrary structured data as output.

3.6.3 OPA and external data sources

In order to get data from external sources, OPA provides several options that can be chosen based on size and frequency of update [?]. For the use case of this work, external data to be pulled is the **scheduling decision** from the GreenOps Scheduler in a synchronous way. The most suitable option is to **pull data during policy evaluation** since scheduling decisions must be made in real-time. The data is pulled from the GreenOps Scheduler (effectively a HTTP server) with the use of the OPA built-in function “*http.send()*” which can be used to send HTTP requests to arbitrary endpoints just like a “curl” command in a shell script [?]. The listings 3.11 and 3.12 respectively show an example of a HTTP request to the GreenOps Scheduler and how the OPA external data pull is implemented in Rego.

```

1 #!/bin/bash
2 curl -X POST -H "Content-Type: application/json" \
3   -d '{
4     "number_of_jobs" : 1,
5     "eligible_regions": ["FR", "IT", "FI"],
6     "deadline": "2025-06-09T10:00:00",
7     "duration": 6,
8     "cpu": 4,
9     "memory": 8,
10    "req_timeout": 180
11  }' http://greenops-scheduler.greenops-scheduler-system.svc.cluster.local/scheduling
12
13 # Response:
14 # {"scheduling_location": "FR", "scheduling_time": "2025-02-06T17:42:11Z"}'

```

Listing 3.11: GreenOps scheduler HTTP request example

```

1 scheduler_url := opa.runtime().env.SCHEDULER_URL
2
3 scheduling_details := http.send({
4   "method": "POST",
5   "url": scheduler_url,
6   "body": {
7     "number_of_jobs": 1, # currently only one job (workload (VM)) can be scheduled
8     # at a time
9     "eligible_regions": eligible_electricity_maps_regions,
10    "deadline": deadline,
11    "duration": duration,
12    "cpu": cpu,
13    "memory": memory,
14    "req_timeout": 10 # seconds, scheduler wants to know the timeout to tune the
      # execution time of the optimization
15  },

```

```

15     "timeout": "10s",
16     "headers": {
17       "Content-Type": "application/json"
18     },
19     "max_retry_attempts": 3,
20   })

```

Listing 3.12: OPA external data pull

3.6.4 OPA integration with Kubernetes

In Kubernetes admission control, policy enforcement is handled by the **Kubernetes API server** itself. OPA makes the policy decisions when queried by the admission controller, but the actual enforcement (namely allowing or denying requests) is executed by Kubernetes' built-in admission control mechanisms. This is effectively one of the core design principles of OPA: to provide policy decisions to external systems, which then enforce the decisions based on their own logic [?]. This workflow is represented in figure 3.9 where **AdmissionReview request** and **AdmissionReview response** are respectively input and output of the whole OPA section. The API Server sends the entire Kubernetes object in the webhook request to OPA. The Kubernetes API server will use the received AdmissionReview response for its decision.

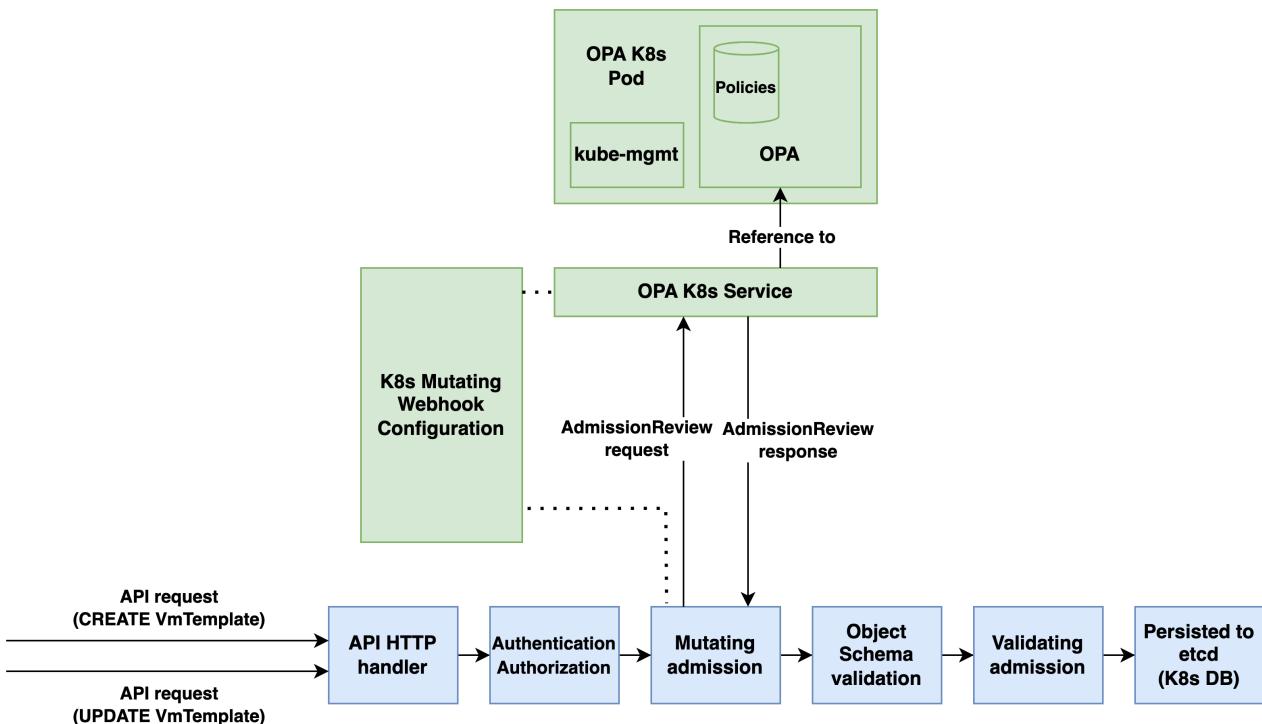


Figure 3.9: Kubernetes mutating webhook and OPA integration

In a Kubernetes deployment, an **OPA server Pod** typically consists of the following containers:

- OPA server container
- **kube-mgmt** container

The kube-mgmt container functions as a **sidecar container** within a Kubernetes Pod. The sidecar container pattern is a common Kubernetes design paradigm in which auxiliary containers run alongside the main application container within the same Pod. These additional containers serve to enhance, extend, or support the primary application's functionality without modifying its core logic [?]. The primary responsibility of kube-mgmt is to replicate Kubernetes resources into the OPA instance (OPA container). This operation is essential for OPA to access and evaluate policies based on real-time cluster state, enabling dynamic policy enforcement. By synchronizing these resources,

kube-mgmt ensures that OPA has an up-to-date view of relevant Kubernetes objects. This is especially useful to enforce policies that deals with naming conflicts, where OPA needs to check existing names in the cluster for the decision [?]. Additionally, it allows for loading policies directly from the Kubernetes cluster by retrieving them in the form of ConfigMaps. This feature is particularly useful when policies need to be dynamically updated based on the current state of the cluster [?]. However, in the system described in this thesis, this latter feature is not employed in the current implementation as we are using policy bundles for policy distribution.

In the current system configuration, the kube-mgmt container is deployed to facilitate resource replication, ensuring that Kubernetes resources, namely VmTemplate resources, are synchronized with the OPA instance. However, at present, no policy requires interrogation of VmTemplate resources that are already present in the system. Looking ahead, future policies could leverage VmTemplate resource information to enforce naming conflict resolution, quota management, or additional constraints.

3.6.5 OPA policies

As OPA official documentation describes, when the Kubernetes AdmissionReview request from the webhook arrives, it is binded to the OPA input document and generates the default, “root”, decision: *system.main*

The root policy, in the case of Kubernetes admission control, is responsible for generating the AdmissionReview response in accordance with the Kubernetes API specifications. It is the duty of the policy developer to write Rego code that produces a well-formed AdmissionReview response, ensuring that the OPA server can then correctly communicate its decision to the Kubernetes admission controller.

It is deemed useful to show one of the simplest and common example of a OPA policy in the **Kubernetes admission control context**. That is: to ensure all images for Kubernetes Pods come from a trusted registry, namely *unitn.it*.

It is important to note that, in this case, due to the simplicity of the policy, no additional contextual data in JSON format is required.

policy compilation policy are compiled compile time errors like merge errors if data is clashing for instance

```

1 deny contains msg if {
2     input.request.kind.kind == "Pod"
3     image := input.request.object.spec.containers[_].image
4     not startswith(image, "unitn.it/")
5     msg := sprintf("image '%v' comes from untrusted registry", [image])
6 }
```

Listing 3.13: Rego policy for Pods registry

```

1 package system
2
3 import data.kubernetes.admission
4
5 main := {
6     "apiVersion": "admission.k8s.io/v1",
7     "kind": "AdmissionReview",
8     "response": response,
9 }
10
11 default uid := ""
12
13 uid := input.request.uid
14
15 response := {
16     "allowed": false,
17     "uid": uid,
18     "status": {"message": reason},
19 } if {
20     reason := concat(", ", admission.deny)
```

```

21     reason != ""
22 }
23
24 else := {"allowed": true, "uid": uid}

```

Listing 3.14: Rego “root” policy (system.main)

```

1 {
2     "apiVersion": "admission.k8s.io/v1",
3     "kind": "AdmissionReview",
4     "request": {
5         "kind": {
6             "group": "",
7             "kind": "Pod",
8             "version": "v1"
9         },
10        "object": {
11            "metadata": {
12                "name": "myapp"
13            },
14            "spec": {
15                "containers": [
16                    {
17                        "image": "bitnami/node:22",
18                        "name": "nodejs"
19                    }
20                ]
21            }
22        }
23    }
24 }

```

Listing 3.15: AdmissionReview request

```

1 {
2     "apiVersion": "admission.k8s.io/v1",
3     "kind": "AdmissionReview",
4     "response": {
5         "allowed": false
6         "status": {
7             "message": "image 'bitnami/node:22' comes from untrusted
8                 registry"
9         }
10    }
}

```

Listing 3.16: AdmissionReview response

Therefore, in this specific case, the creation of the Kubernetes Pod will be **denied**. OPA is responsible for **decision-making**, determining that the request do not complies with the defined policies, while the Kubernetes API server, using the AdmissionReview response generated by OPA, handles **policy enforcement**, effectively rejecting the CREATE request since it violates the specified rules.

3.6.6 OPA policy bundles

An OPA policy bundle is a collection of policies and optional associated contextual data. More precisely, a bundle is a standardized way to package policies, facilitating version control and distribution [?]. As a matter of fact, a single policy bundle can be potentially used by multiple OPA instances. A policy bundle mainly consists of:

- **Rego policy files** defining the logic.

- **Data files** (in JSON or YAML format) containing contextual information required for policy evaluation (e.g., cloud region mappings).

Policy bundles can be distributed through a variety of mechanisms such as remote HTTP servers (e.g., NGINX) and object storage services (e.g., Amazon S3, Google Cloud Storage, Azure Blob Storage) [?]. One of the most convenient approaches is packaging them as **OCI (Open Container Initiative) images** [?] and this is the approach adopted in the system described in this thesis.

Once packaged as OCI images, policy bundles can be pulled by OPA servers from a container registry at predefined time intervals. This allows policy updates to be deployed in OPA **without requiring manual intervention or service restarts**, ensuring that enforcement mechanisms remain up to date with the latest compliance requirements identified and implemented by the organization. This is crucial for instance when dealing with **critical security policies** that need to be updated frequently, maybe in response to the discoveries of new CVEs. In the context of our system, such timely updates are not essential but the OPA is designed to be able to handle them if needed. To ensure continuous policy enforcement while maintaining high operational efficiency, a CI/CD approach is adopted for policy management in the context of our system. As a matter of fact, policies are maintained in a **version-controlled hosted repository** (i.e., on GitHub), where updates like tagging (“git tag”) trigger an automated pipeline (e.g., using GitHub Actions) responsible for building, packaging into a OCI image, and publishing the policy bundle to a container registry (e.g., Docker Hub). One of the major advantages of this approach is the ability to dynamically update policies without requiring OPA pods to restart as there is an **hot-reload** of policies done at application level by OPA (“loaded on the fly”) [?]. This is particularly useful in production environments where service availability is critical and downtime must be minimized. The overall process of policy distribution is illustrated in figure 3.10.

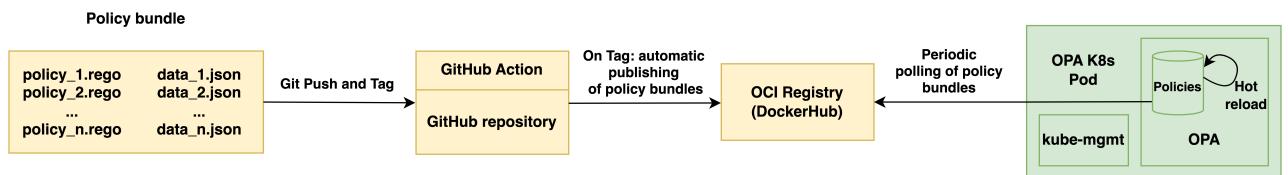


Figure 3.10: OPA policy bundles

By leveraging OCI images for policy distribution and implementing a fully automated CI/CD pipeline, our system ensures that policy enforcement remains consistent, up to date, and highly available across all OPA instances. This approach aligns with modern DevOps practices, enabling organizations to maintain a high level of security and compliance without compromising operational efficiency.

3.6.7 Latency policy

A representative example of a policy aligned with Service Level Objectives (SLOs) or Service Level Agreements (SLAs) is the latency policy described in this section. Given an **origin region** and a **maximum latency threshold** (expressed in milliseconds), the objective is to determine a **set of eligible regions** where the inter-regional latency between the origin and each region in the set is equal to or below the specified threshold. Enforcing such constraints helps mitigate the so-called “**black hole phenomenon**” in the GreenOps use case, where all virtual machines (VMs) would otherwise be scheduled in a region with generally low carbon intensity, without considering additional constraints or performance requirements. By incorporating similar performance-aware policies, organizations can achieve a balance between environmental impact, performance, and service reliability. The proposed flexible system enables organizations to fine-tune these factors according to their specific requirements or those of their users. This policy demonstrates the flexibility of OPA in handling diverse compliance scenarios. It is the responsibility of the policy developer to design an appropriate strategy for encoding relevant information into **well-structured JSON data models**, e.g., a latency matrix. Proper structuring ensures efficient policy evaluation, maintainability and extendability.

Figure 3.11 illustrates a small example (4 regions subset) of a latency matrix, where each cell represents the latency between two regions. The matrix can be encoded in JSON format as illustrated in listing 3.17, allowing for easy integration with OPA policies. The “Latency policy” then uses this matrix to determine eligible regions based on the origin region and maximum latency threshold.

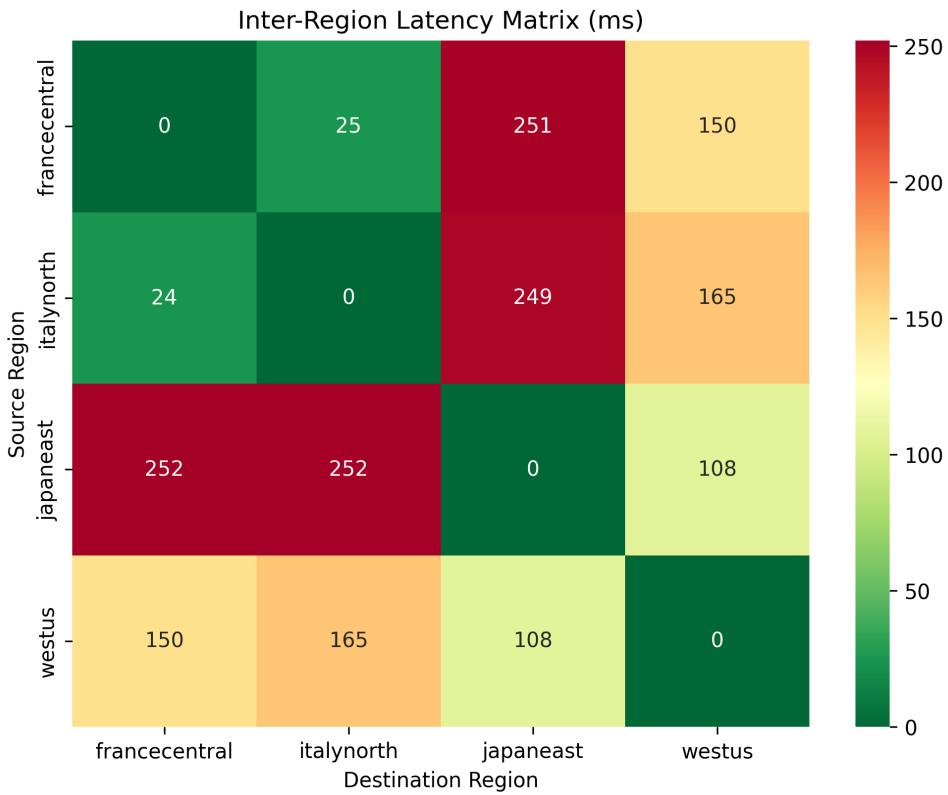


Figure 3.11: Latency matrix example (Azure regions subset)

```

1  {
2    "italynorth": {
3      "italynorth": 0,
4      "japaneast": 249,
5      "francecentral": 24,
6      "westus": 165
7    },
8    "japaneast": {
9      "italynorth": 252,
10     "japaneast": 0,
11     "francecentral": 252,
12     "westus": 108
13   },
14   "francecentral": {
15     "italynorth": 25,
16     "japaneast": 251,
17     "francecentral": 0,
18     "westus": 150
19   },
20   "westus": {
21     "italynorth": 165,
22     "japaneast": 108,
23     "francecentral": 150,
24     "westus": 0
25   }
26 }
```

Listing 3.17: Latency matrix example encoded in JSON format

For what concern data sources, cloud-specific inter-regional latency data were obtained in different ways for each cloud provider. Microsoft Azure provides monthly Percentile P50 round trip times between Azure regions in its documentation [?]. This data was scraped, merged and used to build the latency matrix for Azure regions. For AWS a similar approach was used, but the data was obtained from a third-party website that provides latency data between AWS regions calculating using AWS Lambda functions [?]. For Google Cloud Platform, no official data was found, so the latency matrix was built using synthetic data.

3.6.8 GDPR policy

Another policy configured in the system is the “GDPR Policy”, which ensures that virtual machines (VMs) are deployed in cloud regions that reside in countries of the European Union. The policy is based on the principle of **set intersection**. One set consists of the eligible regions determined by other constraints, such as latency requirements. The other set includes cloud provider regions that are physically located within European Union (EU) countries. The intersection of these two sets defines the final list of allowed deployment regions, restricting workloads to EU-based data centers. Since each cloud provider has its own regional distribution, the list of EU-compliant regions is provider-specific and is encoded as contextual data in JSON format. This allows for flexibility and easy updates when cloud providers introduce new regions.

It must be noted that this policy is **not intended to be a comprehensive GDPR compliance solution**, but rather a basic example of how OPA can enforce **data residency requirements in a multi-cloud environment**. Organizations with more stringent GDPR compliance needs should consider additional measures.

3.6.9 Scheduling outcome policy

In the context of this work, the “Scheduling outcome policy” is a policy that determines the scheduling decision for a given workload based on the output of the GreenOps Scheduler which is queried in real-time, as described in section 3.6.3. The inputs and outputs of this policy are as follows:

- Given:
 - {CPU, RAM, duration, deadline, max latency}: set at request time
 - eligible cloud providers: set at request time or in policies
 - origin region: set in policies
 - GDPR compliancy: set in policies
 - inter-region latency matrix: stored in policy data
- Output decision:
 - provider
 - schedulingLocation
 - schedulingTime

It is therefore a policy that has the duty to mutate (patch) the VmTemplate Kubernetes resource (Generic VM), adding the scheduling information (provider, schedulingLocation, schedulingTime) to the resource. According to Kubernetes documentation, this can be done using “*patch*” and “*patchType*” fields in the AdmissionReview response [?]. The “*patchType*” field must be “*JSONPatch*” and the “*patch*” field must contain a base64-encoded array of JSON patch operations to be applied to the resource. JSON Patch is a format for describing changes to a JSON document which avoid the need to send the entire document when only a part of it has changed. Effectively, only deltas are sent back to the requester which are themselves JSON documents. The format is defined in RFC 6902 from the IETF [?]. As an example, a single patch operation is the one shown in listing 3.18, where a new field “schedulingTime” is added to the resource.

```

1 # schedulingTime is data coming from the GreenOps Scheduler
2 patchCode = {
3     "op": "add",
4     "path": "/spec/schedulingTime",
5     "value": schedulingTime,
6 }

```

Listing 3.18: JSON Patch example

3.6.10 OPA Data mapping

OPA is flexible enough to handle **data mapping operations** between different data models, enabling seamless integration with external systems. In our GreenOps system, data mapping is essential for translating between ElectricityMaps regions and cloud provider regions. At some point in the system this mapping needed to be done and we deemed that inside the OPA policies was the best place to do it. In particular, this mappings are needed since the GreenOps scheduler knows only the notion of ElectricityMaps regions, and do not possess the knowledge of cloud provider regions. Therefore, a mapping is needed to translate the ElectricityMaps regions to cloud provider regions and vice versa.

The entire data mapping process can be broken down into the following steps:

1. Cloud provider selection (e.g., Azure, AWS, GCP): this determines the set of cloud provider regions.
2. Latency filtering: this step can only be done with cloud provider-specific latencies and determines the eligible regions.
3. GDPR compliance filtering (if enabled): this step ensures that only regions in the EU are selected.
4. ElectricityMaps region mapping: this step maps the eligible cloud provider regions to ElectricityMaps regions.
5. Scheduling outcome: this step determines the scheduling region (ElectricityMaps region) based on the output of the GreenOps Scheduler.
6. The ElectricityMaps region is then mapped back to the cloud provider region as final step.

This process is illustrated in figure 3.12.

The Rego code in listing 3.19 illustrates the Rego functions used for data mapping between cloud provider regions and ElectricityMaps regions.

```

1 # Utility functions to map between cloud provider regions
2 # and ElectricityMaps regions
3
4 map_to_electricitymaps(eligible_regions, provider) = em_regions if {
5     em_regions := {
6         region.ElectricityMapsName |
7             some eligible_region;
8             some region;
9             eligible_region = eligible_regions[_];
10            region = data[provider].cloud_regions[_];
11            region.Name == eligible_region
12            region.ElectricityMapsName != ""
13            region.ElectricityMapsName != "Unknown"
14        }
15    }
16
17 map_from_electricitymaps(em_region, provider) = cloud_region if {
18     some region;
19     region = data[provider].cloud_regions[_];
20     region.ElectricityMapsName == em_region;
21     cloud_region := region.Name
22 }

```

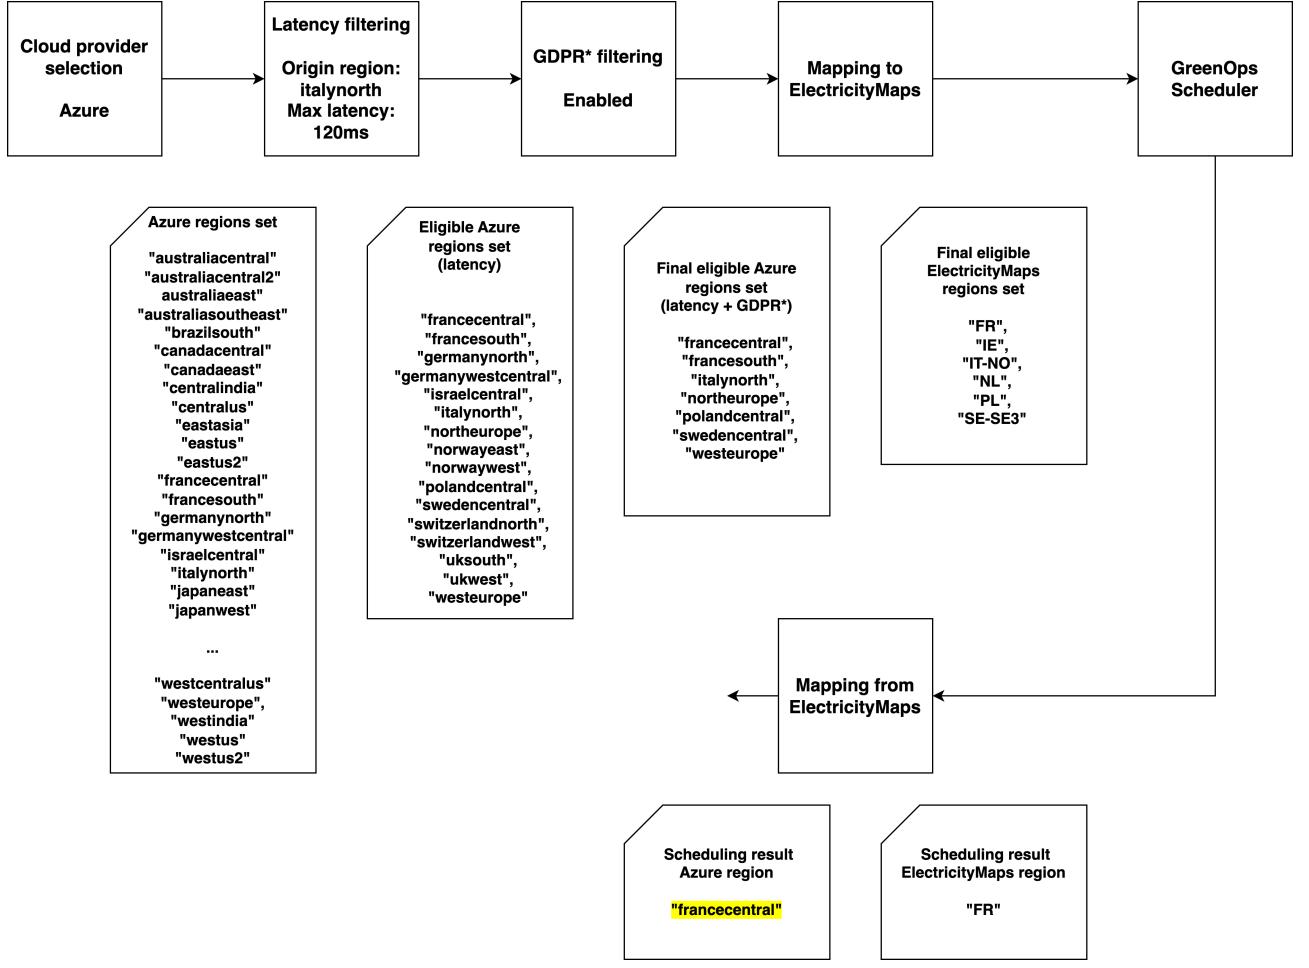


Figure 3.12: OPA Data mapping

Listing 3.19: Rego data mapping

3.6.11 OPA integration in the system

In this section we describe the integration of OPA in the system. The entire architecture of the system and the integration of OPA is illustrated in figure 3.13. OPA has the role of **mutating webhook server** which is consulted by the Kubernetes API server when a CREATE or UPDATE operation is performed on a VmTemplate Custom Resource. The OPA server is responsible for evaluating the policies and returning the AdmissionReview response to the API server. The AdmissionReview response contains the decision of the policy evaluation (i.e. scheduling outcome) and the JSON patch operations to be applied to the VmTemplate resource by the Kubernetes API server. OPA is periodically polling the policy bundles from an external container registry (e.g., DockerHub) to ensure that the policies are up to date. The main policy, namely the “scheduling outcome policy”, is responsible for determining the scheduling decision based on the output of the GreenOps Scheduler called in real-time at request time, as described in section 3.6.9. OPA is also responsible for data mapping operations between ElectricityMaps regions and cloud provider regions, as described in section ???. The system is designed to be highly flexible and extensible, allowing for the addition of new policies and data mappings as needed.

Day 2 operations

We can define as ***Day 2 operations*** all the management operations that are performed after the deployment of a resource. These operations comprise tasks such as scaling up or down a VM

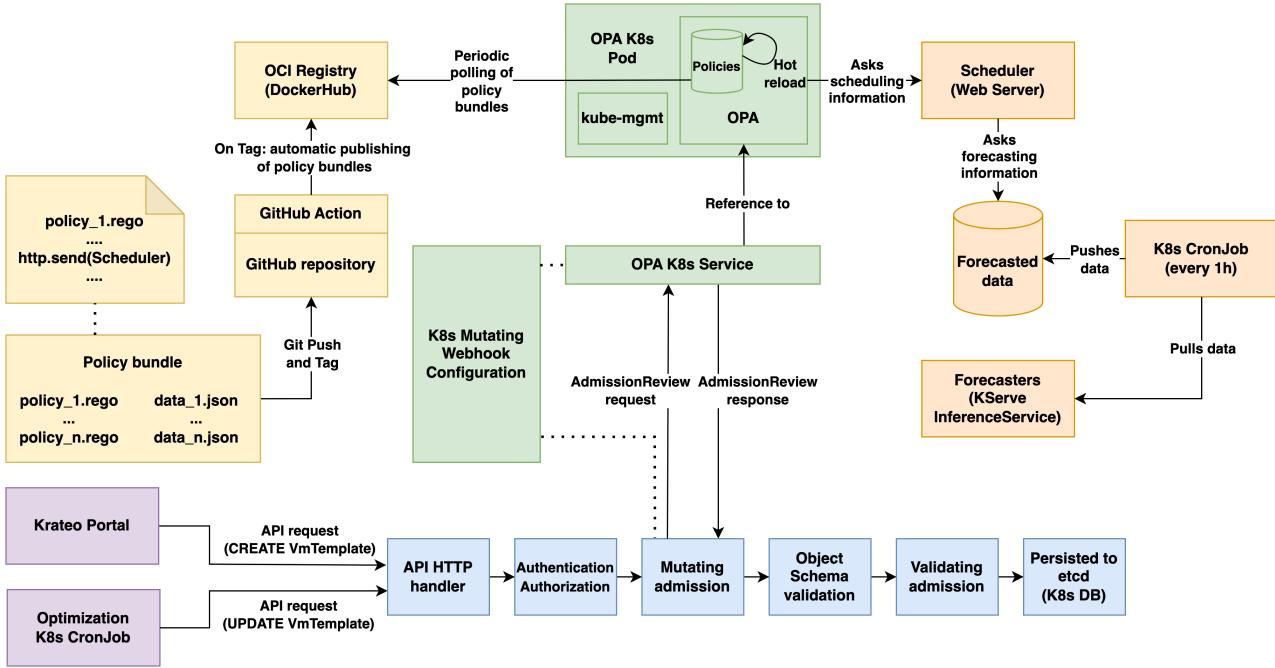


Figure 3.13: General architecture

based on the load, stopping a VM during off-peak hours, or migrating a VM to a different region to optimize costs. As we previously described, the mutating webhook configuration is set on both the **CREATE** and **UPDATE** operations. A possible **UPDATE** operation trigger could be Kubernetes Cronjob that attach a label “*“greenops-optimization”*: “123456789”” to the VmTemplate resource at a specific time of the day. This could be useful to trigger specific policies that are only applied during the day 2 operations. In addition to that, the **UPDATE** operation could be useful for VMs that have already been scheduled in a distant future (due to their deadline very far in the future) to obtain a new scheduling decision based on new conditions (e.g. more recent carbon intensity forecast).

3.6.12 OPA Gatekeeper

OPA Gatekeeper is a Kubernetes-native policy engine that extends OPA with **Custom Resources (CRs)** and controllers to enforce policies across a Kubernetes cluster. It integrates natively with Kubernetes and provides a declarative approach to defining and enforcing policies using Kubernetes Custom Resources (CRs). This makes it an excellent choice for basic and standard policy enforcement scenarios, such as RBAC (Role-Based Access Control), security compliance, and resource constraints. However, while OPA Gatekeeper is well-suited for simple use cases, it presents **limitations** when addressing complex policy requirements, particularly when policies involve **mutations** or require access to **external data sources**. These limitations make it unsuitable for the specific challenges tackled in this system. Therefore, after an initial investigation and Proof of Concept implementation, we decided to use the standard OPA server for policy enforcement mainly due to the flexibility it provides in handling diverse scenarios.

To illustrate the differences between a standard OPA policy and an OPA Gatekeeper policy, we present two examples:

- a simple Rego policy that enforces a basic constraint on Pod creation in a Kubernetes cluster.
- the corresponding policy implemented as an OPA Gatekeeper **ConstraintTemplate** and **Constraint** Kubernetes resources.

The first example demonstrates a standalone Rego policy, which can be evaluated directly by an OPA instance. While this approach is flexible and allows for fine-grained policy definition, it requires manual integration into the system, including policy distribution and enforcement setup.

```

1 package kubernetes.admission
2
3 deny[msg] {

```

```

4   input.request.kind.kind == "Pod"
5   input.request.object.metadata.namespace == "restricted"
6   msg := "Pods cannot be created in the 'restricted' namespace."
7 }
```

Listing 3.20: Simple OPA Rego Policy

The second example, illustrated in listing 3.21 utilizes OPA Gatekeeper, which extends Kubernetes with Kubernetes-native Custom Resource Definitions (CRDs), enabling declarative policy management. By using a ConstraintTemplate, policies can be enforced dynamically through Kubernetes, making them easier to distribute and manage. In other words, with this kind of setting, OPA policy bundles are not employed in the same way as in the standard OPA server. Instead, policies are defined as Kubernetes resources, allowing for more straightforward policy enforcement and management within a Kubernetes environment.

```

1 apiVersion: templates.gatekeeper.sh/v1
2 kind: ConstraintTemplate
3 metadata:
4   name: podnamespaceconstraint
5 spec:
6   crd:
7     spec:
8       names:
9         kind: PodNamespaceConstraint
10      targets:
11        - target: admission.k8s.gatekeeper.sh
12          rego: |
13            package kubernetes.admission
14            deny[msg] {
15              input.review.object.metadata.namespace == "restricted"
16              msg := "Pods cannot be created in the 'restricted' namespace."
17 }
```

Listing 3.21: OPA Gatekeeper ConstraintTemplate

```

1 apiVersion: constraints.gatekeeper.sh/v1beta1
2 kind: PodNamespaceConstraint
3 metadata:
4   name: restrict-namespace
5 spec:
6   match:
7     kinds:
8       - apiGroups: []
9         kinds: ["Pod"]
10        parameters: {}
```

Listing 3.22: OPA Gatekeeper Constraint

In the example, the policy is defined as a ConstraintTemplate, which is then instantiated as a Constraint Custom Resource of kind defined in the ConstraintTemplate. The ConstraintTemplate specifies the Rego policy logic, while the Constraint defines the target resources and parameters for policy enforcement. Therefore a ConstraintTemplate can be used by multiple Constraints, allowing for policy reuse.

OPA Gatekeeper also provides additional Kubernetes Custom Resources called *mutators* (Assign, AssignMetadata, AssignImage, ModifySet) that allow modifying resource fields without writing Rego code. These mutators are useful for simple transformations, such as setting default labels or annotations. However simultaneous mutation of multiple fields leveraging external data is not supported [?]. This limitation, in the context of our system, determined the choice of the standard OPA server for policy enforcement.

It must be noted that OPA Gatekeeper limitations could be potentially addressed in future releases, making it a more viable option for complex policy enforcement scenarios. However, for the current

system requirements, the standard OPA server was deemed more suitable due to its flexibility.

3.6.13 OPA advanced features

It is deemed useful to mention some of the advanced features of OPA that were not employed in the first iteration of the system described in this thesis but could be potentially useful in future developments or in other contexts where OPA is used.

To ensure data integrity of policies (i.e., to prevent unauthorized modifications or MITM (Man in the Middle) attacks) and to provide a mechanism for verifying the authenticity of policies, OPA provides a **policy signing** feature [?]. This feature allow policy developer to digitally sign their policies bundles by adding a file named “*.signature.json*”.

OPA also provides a more efficient way to distribute policies using **Delta Bundles** [?]. Normally, when a policy bundle is downloaded, OPA will download the entire bundle, erase and overwrite the current policies and data with the new ones. Delta Bundles are composed of a single “patch.json” file that contains a set of JSON Patch operations that can be applied to the current data to update it to the new version. Currently only data updates are supported (data.json) and not policy updates (policy.rego) [?].

3.7 MLOps infrastructure

MLOps is the abbreviation of “**Machine Learning Operations**”, and it broadly refers to a set of methods designed to improve workflow procedures and automate machine learning deployments. It enables the reliable and efficient management, maintenance and deployment of models at scale [?]. A MLOps infrastructure is not necessarily required for multi-cloud resource management, but it is believed that AI models will be utilized in the future more and more to get **scheduling and management decisions**, as evidenced by recent studies discussed in section 2.4.3. It is therefore deemed important to describe the MLOps infrastructure deployed in a Kubernetes environment and leveraged by the system described in this thesis.

3.7.1 MLOps purpose

In a way, MLOps implements DevOps principles, tools and practices into typical Machine Learning workflows. Its main purpose is to effectively industrialize the machine learning models lifecycle, enabling faster model development, selection, and deployment to production compared to traditional manual approaches. Some principles of MLOps can be summarized as follows [?]:

- **Automation:** automate the entire model lifecycle, from training to deployment.
- **Versioning:** track and version models (with related data and code).
- **Reproducibility:** ensure that various model versions can be reproduced at any time.
- **Monitoring:** monitor models in production to ensure they are performing as expected.
- **Scalability:** scale models to handle increased workloads in a seamless manner.
- **Collaboration:** enable collaboration between data scientists, data engineers, and operations teams.

In the context of the proposed system, the MLOps infrastructure is used to deploy and manage the forecasting models that predict the carbon intensity of the electricity grid in different world regions. In particular, **MLflow** is used for model tracking, model selection, and model storage, while **KServe** is used for model deployment.

3.7.2 MLflow

MLflow is an open-source platform designed to facilitate and enhance the overall machine learning lifecycle. In particular, it offers tooling for **experiment tracking**, **model management**, and **model storage**. It is compatible with various ML frameworks, including scikit-learn, PyTorch, TensorFlow, and XGBoost. In particular, in the case of our system, PyTorch is the ML framework used for the forecasting models making MLflow a suitable choice for model management. For what concerns model deployment, MLflow does not provide a built-in solution, but it is compatible with other tools like KServe which will be described in section 3.7.3.

MLflow Tracking Server

MLflow Tracking Server is the central component of the MLflow platform which is responsible for logging and storing model training data, parameters, and metrics. It enables ML practitioners to track experiments, compare results, and reproduce models easily in a collaborative environment. MLflow Tracking Server revolves around the concepts of **experiments**, which are collections of **runs**. Each run represents a single model training session, and it contains information such as parameters, metrics, and artifacts (e.g., model files). Selected models can then be registered in the MLflow Model Registry, which is essentially an optional subset of the selected models that are ready for deployment. Model selection can be done with the aid of a user interface that allows ML practitioners to visualize and compare the results of different experiments, making it easier to select the best model for deployment. MLflow Tracking Server, as the name suggests, is a server that is constantly waiting for new data to be logged. Said data comes from the various training scripts that are executed by the data scientists or machine learning engineers which are run in training environments. In particular, the training scripts must be instrumented with the **MLflow API calls** to log the data in the remote MLflow Tracking Server. This setting is very flexible since the training scripts can be run in any environment (e.g., local environment, Universities’ HPC clusters), as long as they have access to the MLflow Tracking Server. It

is deemed important to mention some of the MLflow API calls that are used in the training scripts: the *infer_signature* and *autolog* functions. The *infer_signature* function captures the **input and output schema of the model**, which is important for the deployment of the model. The *autolog* function is used to automatically log the parameters and metrics of the model training session, without the need to manually log them. In particular, each supported ML framework has its own autolog function which is used to automatically log the parameters and metrics of the model training session.

The final phase of a model training session is the automatic creation of a **self-contained directory**, named after the experiment and run IDs, that contains all the necessary files to deploy the model. In particular, the folder contains: the serialized model with model weights, the configuration files (conda.yaml, requirements.txt), and the MLmodel file that contains additional configuration, among which, the model signature. The self-contained folder is then automatically uploaded to the MLflow Tracking Server as an artifact. The following is an example of the structure of the self-contained directory created by MLflow after a PyTorch model training session:

```
model/
├── MLmodel
├── conda.yaml
├── python_env.yaml
└── requirements.txt
data/
└── model.pth
    └── pickle_module.info.txt
```

MLflow deployment configuration

In the context of our system, MLflow is deployed on a Kubernetes cluster with a loosely coupled architecture as can be seen in Figure 3.14, where the MLflow Tracking Server is decoupled from its storage: the metadata store (backend store) and the artifact store. This configuration is the most common in production environments, as it allows for better scalability and environment flexibility. The metadata store chosen for the system is **CrateDB**, while the artifact store is the **SeaweedFS** object storage service.

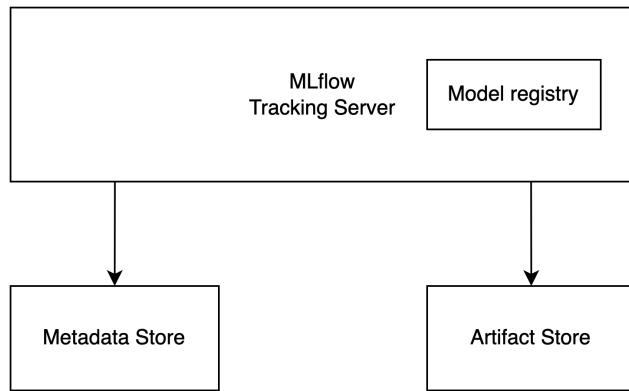


Figure 3.14: MLflow deployment configuration

During the design phase, **alternative configurations** were considered. One of the alternative configurations was to use a single CrateDB instance as both the metadata store and the artifact store. This configuration was not implemented due to the lack of native support of object storage in CrateDB. A second theorized approach was to use a sidecar container with the duty of watching the MLflow Tracking Server local directory in the container filesystem (e.g., using *watchdog* Python package), packaging the model artifacts as an OCI image, and uploading them to an image registry. This approach could be implemented to offer an alternative to the MLflow Tracking Server artifact

store in environment where adding a new storage service is not feasible.

3.7.3 KServe

KServe is an open-source model inference platform that extends Kubernetes with a set of Custom Resource Definitions (CRDs) to **deploy** and scale machine learning models in production environments [?]. KServe can be deployed in several ways, one of which is the “Serverless mode” which is built on top of Istio and Knative, leveraging their powerful capabilities such as automatic scaling. There are two main CRs that can be used to set up a model serving environment: **ServingRuntimes** (or **ClusterServingRuntimes**) and **InferenceServices**. The former are abstractions that define model serving environments, specifying the templates for Kubernetes Pods capable of serving particular model formats. The latter are actually leveraging the available ServingRuntimes to deploy the models in the system. KServe provides several out-of-the-box ClusterServingRuntimes for common model formats, such as TensorFlow, PyTorch, and XGBoost, which can be used to deploy models without the need to define and configure the runtimes themselves. In particular, in our specific use case, since the forecasting models are PyTorch models, packaged by MLflow as described in the previous section, InferenceServices with “kserve-mlserver” ClusterServingRuntime are used. As a matter of fact, this runtime is the one that supports models packaged with MLflow. Listing 3.23 shows an example of the InferenceService Custom Resource used to deploy a model in the system.

```
1 apiVersion: "serving.kserve.io/v1beta1"
2 kind: "InferenceService"
3 metadata:
4   name: "forecaster-IT-NO"
5   namespace: "model-inference"
6 spec:
7   predictor:
8     serviceAccountName: sa-s3creds
9   model:
10    modelFormat:
11      name: mlflow
12    protocolVersion: v2
13    storageUri: s3://mlartifacts/forecaster-IT-NO
```

Listing 3.23: InferenceService Custom Resource example

We can see that the InferenceService CR is quite simple and self-explanatory. The most important fields are the *model* field, which specifies the model format, the protocol version, and the storage URI. The storage URI is the location of the model artifacts in the artifact store, which in the case of the system is the SeaweedFS object storage service with S3 compatibility. At the location specified by the storage URI, the model artifacts are stored in the self-contained directory created by MLflow after the model training session as described in section 3.7.2.

Open Inference Protocol

Interoperability is key in a fast-moving environment as the one of machine learning and AI. Therefore KServe has introduced the **Open Inference Protocol specification** to standardize the communication between inference servers and clients. The Open Inference Protocol has been adopted by several inference servers, including NVIDIA’s Triton and Seldon MLserver. As mentioned in the previous section, the InferenceService CRs used in the system are leveraging the “kserve-mlserver” ClusterServingRuntime, which under the hood uses MLserver that is compliant with the Open Inference Protocol. The list of API endpoints specified by the Open Inference Protocol is shown in Table 3.3.

Model deployment

Our specific use case requires the deployment of multiple models, each corresponding to a different region. This is true if we want to achieve better model performance compared to a single model that tries to predict the carbon intensity of all the regions. The current strategy adopted for the system

API	Verb	Path
Inference	POST	v2/models/[<model_name>]/[<model_version>]/infer
Model Ready	GET	v2/models/<model_name>/[<model_version>]/ready
Model Metadata	GET	v2/models/<model_name>/[<model_version>]/metadata
Server Ready	GET	v2/health/ready
Server Live	GET	v2/health/live
Server Metadata	GET	v2/metadata

Table 3.3: Open Inference Protocol API endpoints specification

is the following: one model per region is deployed, and a generic model is used as a fallback if the specific model is not available.

This translates into the deployment of multiple InferenceServices, each corresponding to a specific region, and one InferenceService that acts as a fallback. This setting introduces a quite large amount of overhead in terms of resources, since each InferenceService set up a whole new set of resources (e.g., large Kubernetes Pods with underlying model serving environments) for each model. Figure 3.15 illustrates the configuration of the InferenceServices used to deploy the forecasting models in the system.

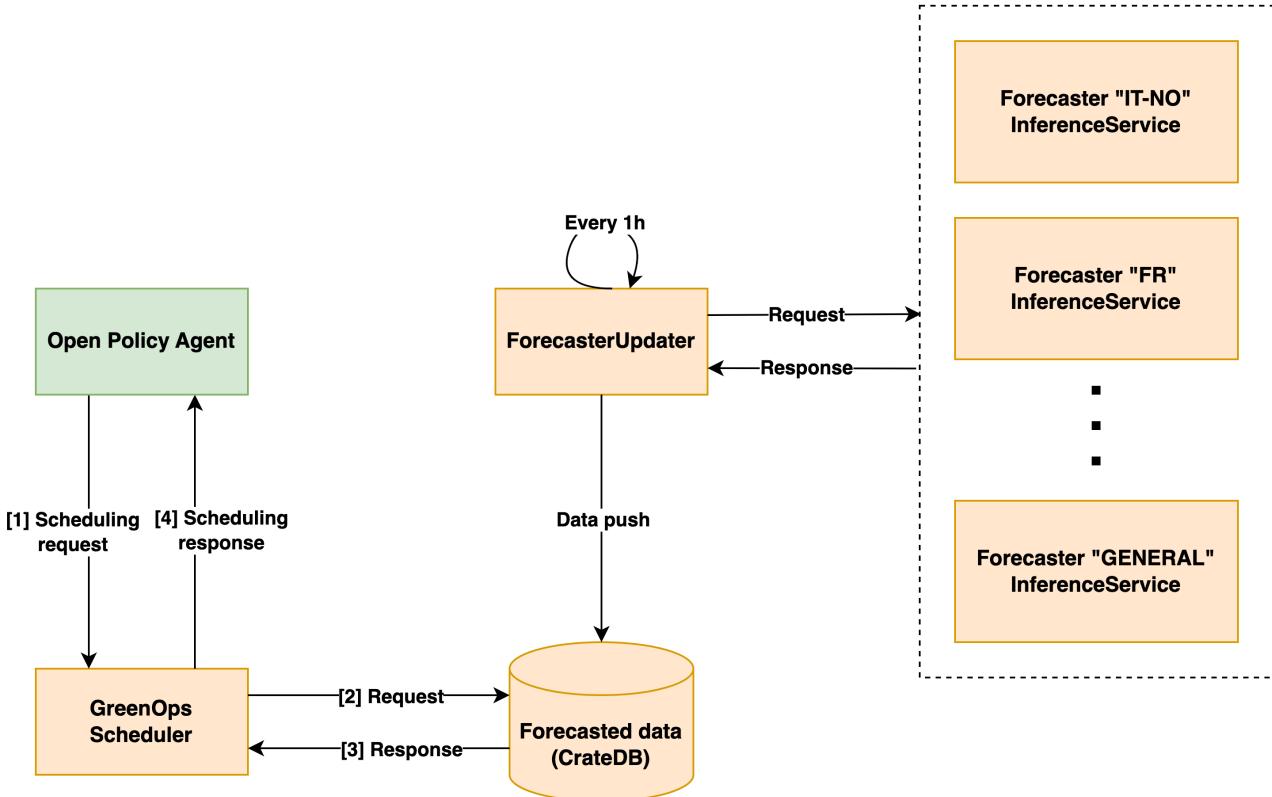


Figure 3.15: Model deployment configuration and GreenOps system integration

Model compatibility

In scenarios where ML models produce outputs in formats not directly compatible with the serving runtime, some workarounds are necessary, such as model wrapping. This process involves adapting the model's input and output interfaces to align with the expected formats of the serving runtime, ensuring a correct model deployment. For instance, in the case of the forecasting models, the output format was not directly compatible with the serving runtime: the output was a custom defined class specific to the model, instead of a single tensor. Therefore, a simple model wrapping was needed to make the model compatible with the serving runtime, effectively operating a selection of output fields

from the model output and returning them as a single tensor.

3.7.4 MLOps general architecture

Figure 3.16 illustrates the general architecture of the MLOps infrastructure deployed in the Kubernetes environment. The architecture consists of two main components, described in the previous sections: MLflow and KServe. In particular, MLflow is used for model tracking, model selection, and model storage, while KServe is leveraged for model deployment.

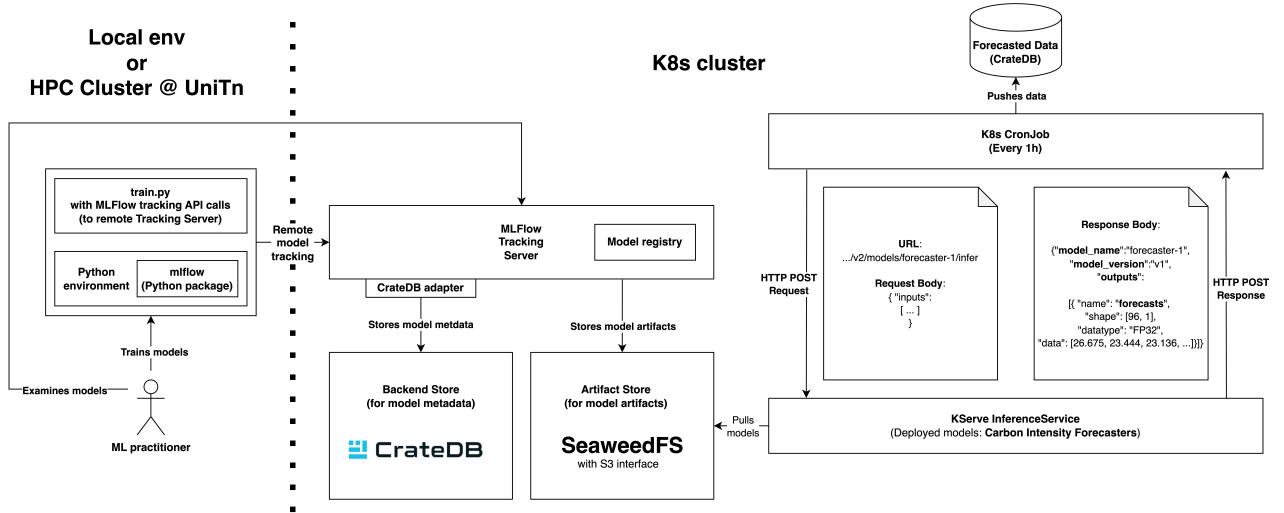


Figure 3.16: MLOps Architecture

3.8 Cloud resource metrics

Metrics are data points that provide information about a system. This information could be related to performance, behavior, resource utilization or any other aspect of the system that could be relevant. In the context of cloud resource management, metrics are essential for monitoring the **performance and health of cloud resources**, enabling organizations to make informed decisions about resource allocation, scaling, and optimization. Focusing on the first use case of the proposed system (i.e., VM scheduling for carbon footprint optimization), carbon metrics could be used to determine the carbon footprint of a scheduled VM. How to measure the carbon footprint of a resource managed in the cloud is a complex and challenging task. The carbon footprint of a cloud resource is influenced by several factors, such as the energy efficiency of the data center where the resource is hosted (PUE), the carbon intensity of the electricity grid in the region where the data center is located, whether the data center uses additional energy sources off the grid, and finally the energy consumption of the resource itself. The last factor is probably the most difficult to measure, as a cloud resource is an **entity that is abstracted from the physical infrastructure**, and therefore it is not straightforward to measure its power consumption. In this first iteration of the system we are dealing for instance with virtual machines (VMs). On the other hand, an on-premises server, for instance, could be potentially equipped with physical sensors that measure the direct power consumption of the server, but this is not the case for any cloud resource, especially since this work is focused on the consumer side of the cloud, where the consumer does not have access to any physical infrastructure. Moreover, public cloud providers do not provide fine grained and real-time data about the carbon footprint of a resource, as they do not adhere yet to a common standard for carbon footprint calculation (e.g., Real Time Cloud proposed by the Green Software Foundation), like they instead do for the FOCUS standard for FinOps. In this section we want to give a brief overview of several metrics and methodologies that can be used to estimate the carbon footprint of a cloud resource. In addition, we will also discuss the **challenges and limitations** that can be encountered when dealing with this kind of task.

3.8.1 System performance metrics

This category of metrics provides information about the health and performance of a system. Usually these metrics are related to bare-metal servers or virtual machines. Other cloud resources that have a higher level of abstraction (e.g., containers, Kubernetes Pods, serverless functions) expose a different set of metrics. This category of metrics are also useful for the “**day 2 operations**”, which are operations that are performed after the resource has been deployed and monitored for a certain period of time, briefly described in section 3.6.11.

System performance metrics comprise, for instance:

- CPU usage
- Memory usage
- Disk usage
- Network usage
- Processes related metrics

A standard strategy to collect these metrics is to use **agents** that are directly installed on the VMs and operate at the Operating System level. Some examples of this kind of agents are the **Prometheus Node exporter** and the **Prometheus Process exporter**. Figure 3.17 shows an example of a standard Prometheus architecture used to collect system performance metrics using the Node exporter and the Process exporter on a set of Linux VMs. In particular, the Node exporter collects metrics such as CPU usage, memory usage, disk usage, and network usage, while the Process exporter collects metrics about the various processes running on the VM (e.g., CPU and memory usage per process group). Metrics are pulled by the Prometheus server, which stores them in a time-series database and makes them available for querying and visualization (e.g., using Grafana). In addition, the Prometheus server can be configured to send alerts based on the collected metrics leveraging the Alertmanager component.

An example of how these metrics can be used for “day 2 operations” is the one put in place by

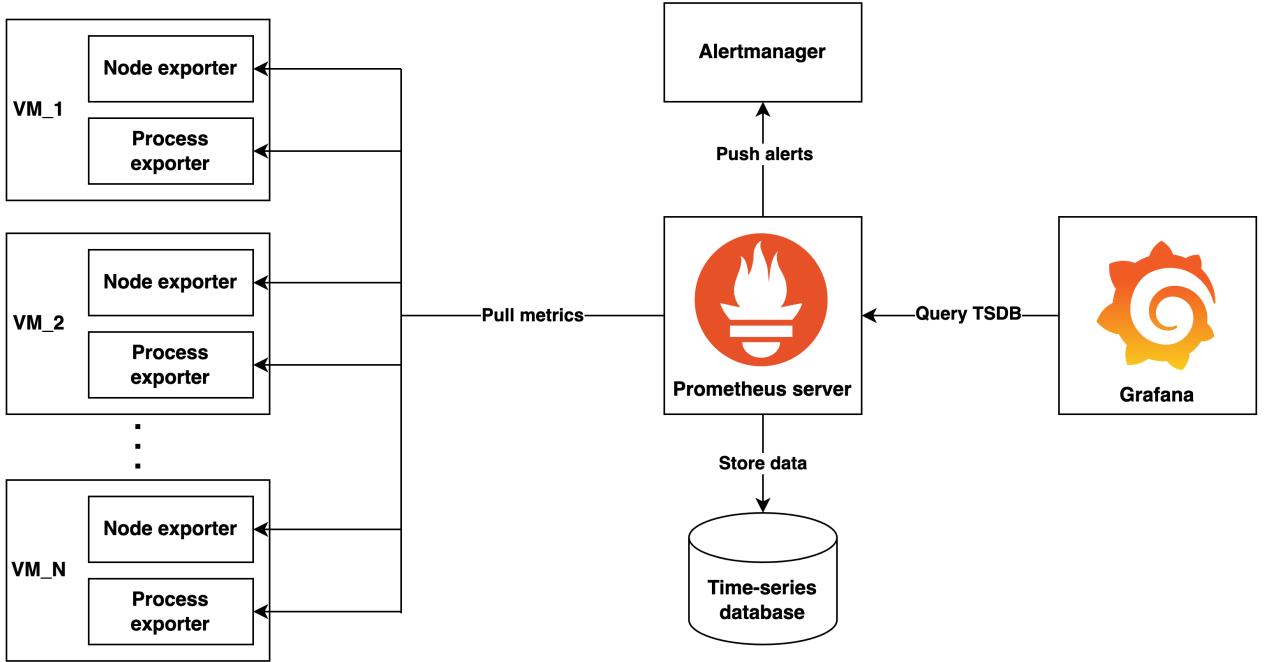


Figure 3.17: Standard Prometheus architecture for system performance metrics collection

the **Krateo Composable FinOps** component [?]. In this case, the system uses a set of Prometheus exporters and scrapers, configured through Kubernetes Custom Resources, to collect metrics about the VMs managed by the system. The optimizations are encoded in a set of Kubernetes Custom Resources (CRs) that are used by the specific operators (e.g., *finops-operator-vm-manager*) to perform operations on the VMs, such as scaling up or down the VMs, stopping them during off-peak hours, etc [?].

As Microsoft Azure documentation describes, there are actually two different categories of metrics that can be collected from a VM: the **guest-level metrics** and the **hypervisor-level or host-level metrics** [?]. The first category of metrics is collected by agents directly installed on the VMs, as described above. The second category of metrics is collected by **cloud provider APIs**, and they provide information about the VMs from the hypervisor or host level. For instance, in the case of Azure Virtual Machines, this data is related to the Hyper-V host that runs the VM [?]. The difference between these two categories of metrics is that the guest-level metrics are more fine-grained and provide information about the actual resource usage of the VM (being collected at the OS level), while the hypervisor-level metrics are more high-level.

To support the collection of metrics, several cloud providers offer agents that can be installed on the provider-specific VMs they manage. One of the goals of these agents compared to standard cloud-agnostic agents is to provide a more convenient experience for the user managing the VMs in the context of the cloud provider. Therefore these agents are usually more integrated with other cloud provider services and APIs. In particular, we can cite the following agents: **Azure Monitor Agent** for Azure VMs, the **Google Ops Agent** for Google Compute Engine instances, and the **Amazon CloudWatch Agent** for AWS EC2 instances.

3.8.2 Power consumption metrics

In the context of our use case, power consumption metrics could be theoretically used to estimate the carbon footprint of a VM instance. However, Public Cloud Providers do not provide real-time data about power consumption of a VM instance.

Scaphandre is an open-source monitoring agent designed to collect energy consumption metrics of a system [?]. Scaphandre represents a promising solution only for on-premises servers, as it relies on the Intel **Running Average Power Limit (RAPL)** sensors to collect power consumption data. As a matter of fact, due to security and privacy concerns, public cloud providers do not expose to tenants

the underlying RAPL sensors that Scaphandre relies on to track energy consumption [?]. Therefore, Scaphandre is unsuitable for our current use case. Another interesting and quite mature project is **Kepler** [?], which is a tool designed to monitor the energy consumption of Kubernetes resources (e.g., Nodes, Pods). Being Kubernetes-centric, Kepler is not suitable for our first use case, as it does not provide real-time data about the power consumption of a generic VM instance. However, Kepler could be potentially used in a future iteration of the system.

3.8.3 Carbon metrics

As briefly mentioned in the beginning of this section, there is no adopted standard adopted by Public Cloud Providers to calculate and expose the carbon footprint of a cloud resource. As described in section 2.5.1, the Green Software Foundation is working on a standard called Real Time Cloud that aims to provide a common standard for carbon footprint calculation. However, this standard is not yet adopted by any major Public Cloud Provider.

Public Cloud Providers provide monthly reports (probably not useful in this case) Export Azure carbon optimization emissions data (Preview) (probably not fine-grained as we want)

Cloud Carbon Footprint Uses cloud provider billing (AWS Cost and Usage Reports with Amazon Athena, GCP Billing Export Table using BigQuery, Azure Consumption Management API). Using these services costs. Electricity Maps API integration is supported (for live grid carbon intensity)

aether calculation engine (<https://aether.green/docs/methodologies/>) only AWS and GCP supported, Azure not yet uses AWS CloudWatch and Google monitoring very small project no live Grid Carbon Intensity Coefficient, they extrapolate data from “governative” data reports <https://aether.green/docs/methodologies/carbon-intensity-coefficient>

A possible approach to estimate the carbon footprint of a VM instance carbond agent ,ust installed on the machine so in this case, probably not useful [?]

Example of a manual approach (not scalable due to tech specs research): <https://devblogs.microsoft.com/sustainable-software/how-can-i-calculate-co2eq-emissions-for-my-azure-vm/>

The critical point here is to get/calculate the energy consumed by a cloud instance, since there are a huge number of technical configurations to find, retrieve and use for calculations.

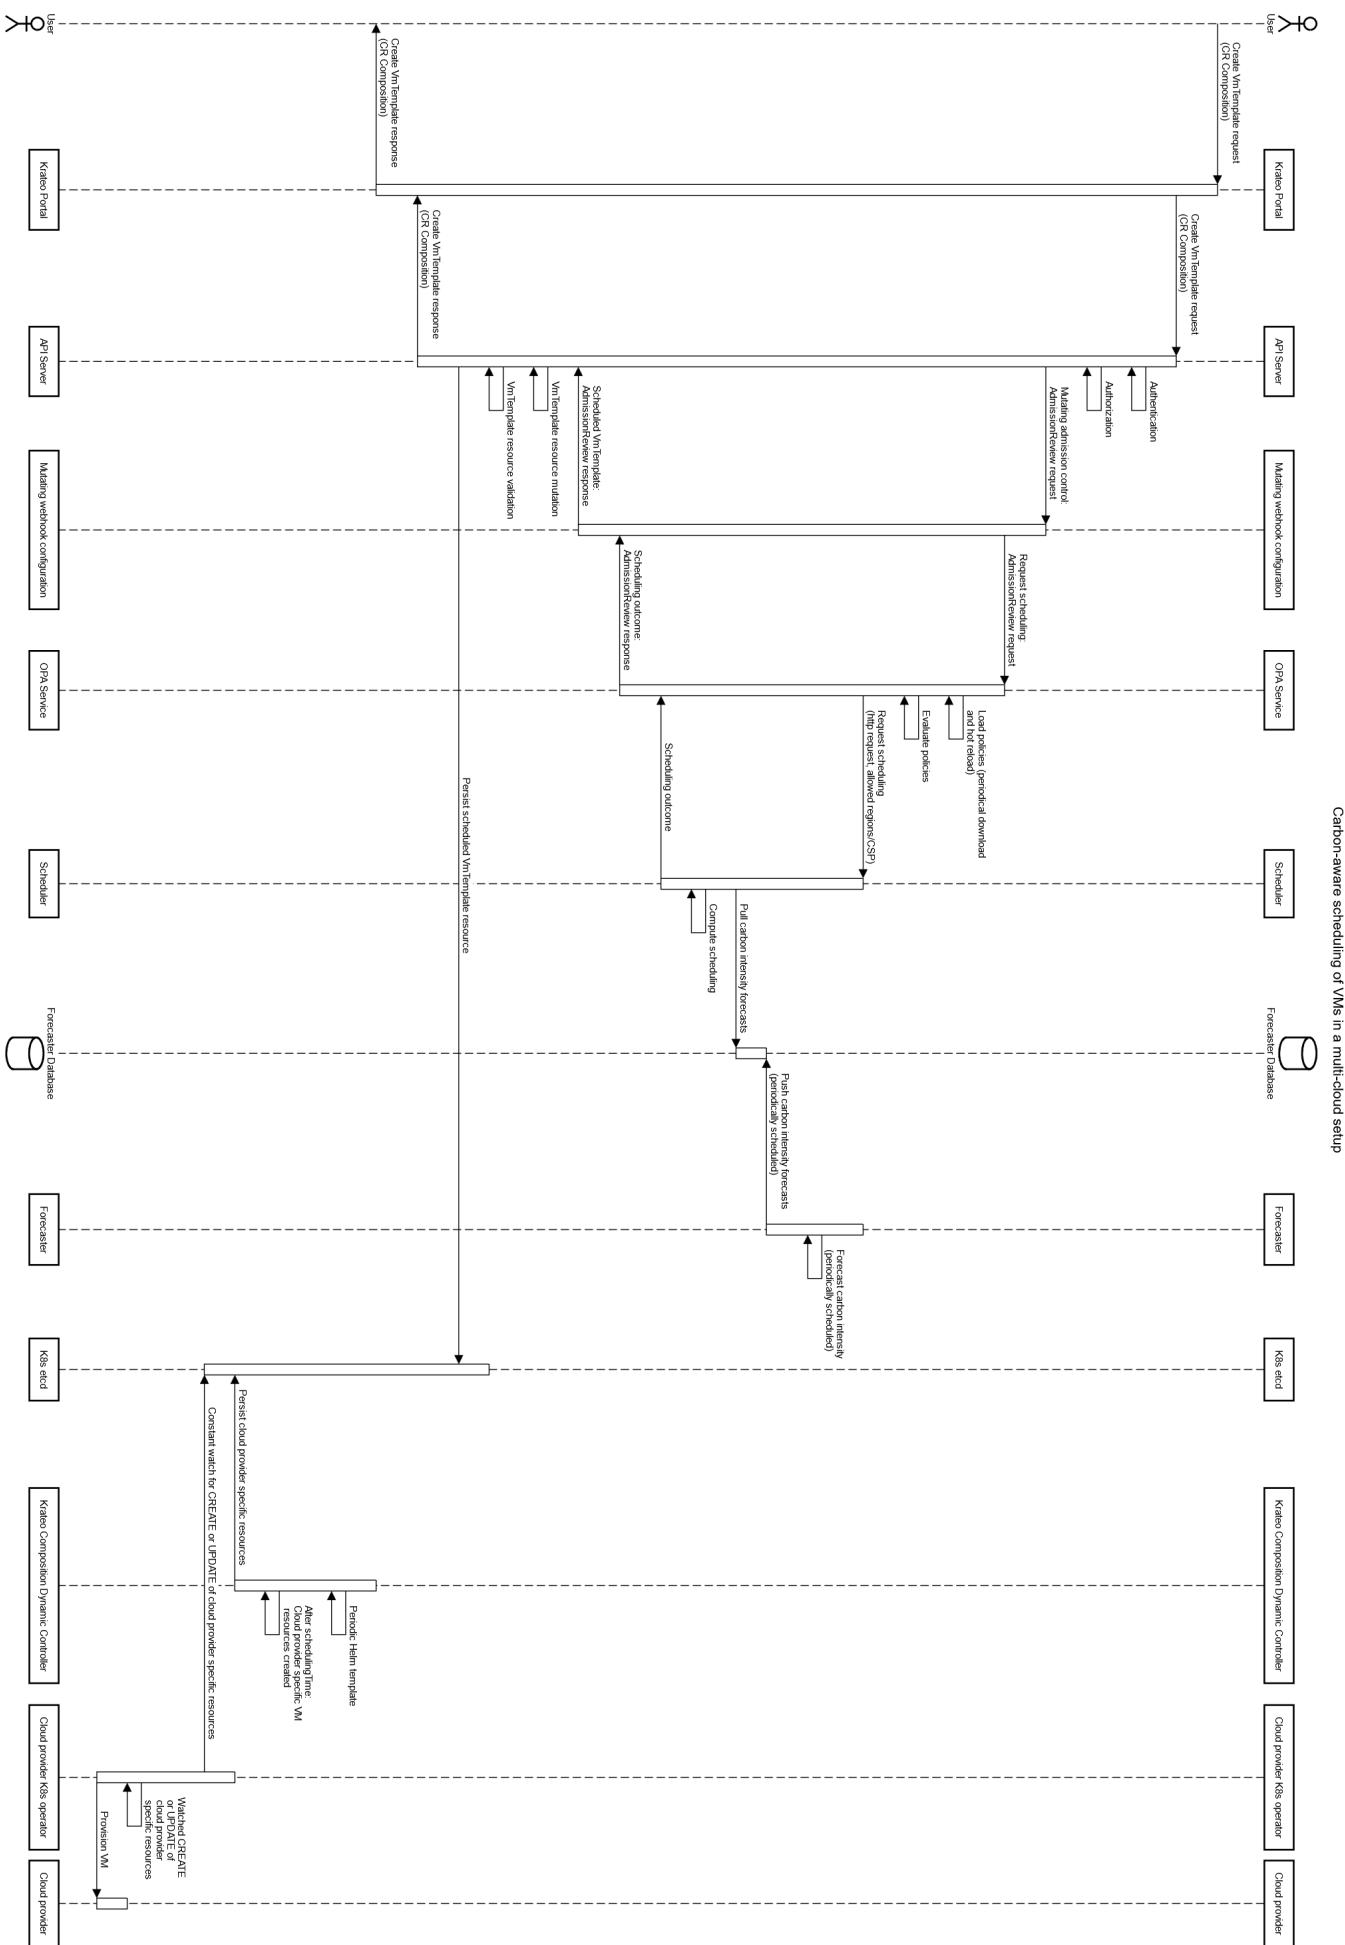
3.8.4 Impact framework potential integration

A potential integration with the Green Software Foundation’s Impact Framework is envisioned for the system. Currently, said integration is not implemented, but it is deemed a valuable addition to the system.

3.9 End-to-End workflow

In this section we will provide a complete end-to-end workflow of the system, from the user request to the final provision of the workload. First, we will describe the workflow in a high-level manner, and then we will provide a detailed sequence diagram (Figure 3.18) that illustrates the interactions between the various components of the system. We must note that the workflow described in this section is a simplified version of the actual system, as it does not include some low level details.

1. User (Developer, Data Scientist, etc) selects the VmTemplate Composition card on the Krateo PlatformOps Portal (UI)
2. User fills in the VmTemplate Composition form with the required fields (e.g., MinCPU, MinRAM, Deadline, Duration, MaxLatency)
3. This triggers and Helm install of the VmTemplate Composition on the Cluster
4. A VmTemplate Composition CREATE API request is sent to the Kubernetes API server
5. Authentication and Authorization checks are performed by the Kubernetes API server
6. The VmTemplate Composition CREATE API request is intercepted by a K8s Mutating Webhook configured with OPA as webhook server
7. AdmissionReview request is sent to the OPA server
8. OPA evaluates the AdmissionReview request against the policies
 - (a) Cloud Provider is selected among the available providers
 - (b) Eligible regions are calculated based on the selected Cloud Provider and the MaxLatency parameter
 - (c) GDPR policy (if enabled) is evaluated and a subset of eligible regions is returned
 - (d) A scheduling request is sent to the GreenOps Scheduler along with the eligible regions and a set of parameters
 - (e) The GreenOps Scheduler returns a decision with the selected region and scheduling time
 - (f) OPA maps the return ElectricityMaps region name to the Cloud Provider region name
 - (g) OPA creates the JSON patch with the provider, schedulingRegion, and schedulingTime fields
 - (h) OPA crafts and sends the AdmissionReview response to the K8s API server
9. The K8s Mutating Webhook receives the AdmissionReview response and mutates the VmTemplate Composition specification
10. K8s API server perform resource validation
11. The VmTemplate Composition is persisted in the etcd database
12. A periodic Helm upgrade operation is done by Krateo Composition Dynamic Controller
13. After schedulingTime is reached, an Helm upgrade operation will install the provider-specific manifests for VM provisioning on the K8s Cluster
14. Cloud provider Operator (e.g., Azure Operator) that is constantly watching for new provider-specific resources is triggered
15. The Cloud provider Operator provisions the VM (and required resources) on the Cloud Provider
16. The VM is up and running on the cloud



4 Conclusion

This chapter concludes the thesis by summarizing the main contributions, and discussing future work.

In particular ... production-ready system

4.1 End-to-end integrated test

final result A comprehensive end-to-end integrated test has been carried out on a Kubernetes cluster
this was used to validate the system

4.2 GreenOps system evaluation

[maybe not really in the scope] [MORE RELATED TO REDI'S THESIS]

4.2.1 Theoretic upper bound

(how close can we get, masachussets amherest group)

4.2.2 Baseline definition

We should prepare one or more baseline scheduling that will be used as a baseline and compared with
a carbon-aware scheduling proposed by our system.

4.2.3 Black hole phenomenon

How to deal with the so-called “Black hole” phenomenon? That is, if 100 workload scheduling arrives
at some point, there is the possibility that the outcome of the system we are building is: “schedule
all workloads in Norway” where Norway is the region with least carbon intensity at that moment.
This phenomenon came up also in a previous meeting but it is not clear if this could be a problem
etc.. A probable differentiator could be the max latency field of the workload request. Other service
requirements could contribute to this as well.

(how it is countered)

4.2.4 Side effects

Maybe out of scope of this work, side effects, big picture. What happens if a big percentage of
companies that relies on cloud services starts to adopt carbon-aware scheduling of their workloads?
We tend to image cloud providers or even cloud regions as an infinite pool of resources, and at a certain
level it is almost like that. But could carbon aware scheduling have larger, not foreseen, side effects?
Is this a responsibility of who schedules? Shall schedulers be responsible for the load on regions? Like
self-imposing some sort of limits/caps.

4.2.5 Preliminary evaluation

4.3 Future improvements

day2 operations we are ready for this

Scaling down a VM (example of Day 2 operations) From: (4 vCPU, 8 GiB RAM) To: (2 vCPU,
4 GiB RAM)

This use case is meaningful for workloads with durations in the order of at least days. Otherwise,
for short-lived workload this use case does not make sense. And in the case of workloads with days as
duration, time and geographical shifting is not that relevant.

This use case will leverage system and performance metrics.

support for other resources we need templates operators

from only vm to other resources

example of a listing of resources that could be supported

4.3.1 Multi model serving

"The original design of KServe deploys one model per InferenceService. But, when dealing with a large number of models, its 'one model, one server' paradigm presents challenges for a Kubernetes cluster."

kserve model mesh instead of several InferenceService there is a lot of overhead in the current configuration

how much is better to use more models instead of one generic model